

# מבוא לראייה ממוחשבת – 22928 2016א

מנחה: אמיר אגוזי  
[egozi5@gmail.com](mailto:egozi5@gmail.com)

מפגש מס' 1

# הנושאים להיום:

- אדמיניסטרציה
- מבוא לקורס
- תמונות
- צבע
- רעש ופילטרים מרחביים
- איתור סף - edge detection

# אדמיניסטרציה

- **מנחה** (ומרכז): אמיר אגוזי, [egozi5@gmail.com](mailto:egozi5@gmail.com)
  - שעות קבלה טלפוניות - שלישי 17:00-19:00, 050-2773676.
- **חומר הקורס**
  - 12 הרצאות של ד"ר טל הסנר (הרצאה בשבוע)
  - בנוסף לחומרים נוספים שיתפרסמו באתר
- **מפגשי הנחיה**
  - כל שבועיים (בדר"כ) = שתי הרצאות/יחידות
  - יש לצפות בהרצאות הרלוונטיות
  - במפגשים - דיון והרחבה בנושאי ההרצאות

# מטלות

- 3 מטלות – ניקוד כל אחת – 5%.
- 2 מטלות תיכנות ב – Python, מטלה 1 – שאלות פתוחות.
- חובה להגיש 2 מתוך 3.
- הגשה דרך מערכת מטלות בלבד.

# פרויקט מסכם - תחרות

- התמודדות עם dataset "אמיתי"
- מימוש של שיטה קיימת
- פיתוח ומימוש רעיונות מקוריים
- הערכה יחסית

# Python

- הוראות התקנה ל-Anaconda באתר.
- עבודה בעיקר עם OpenCV
- חבילות נוספות:
  - Numpy, matplotlib, scikit-learn
  - Scikit-image
  - Scipy.ndimage
  - PIL

# What is computer vision?

The goal of computer vision is to develop algorithms that allow computer to “see”.

Also called:

- Image understanding
- Image analysis
- Machine vision

# General visual perception is hard





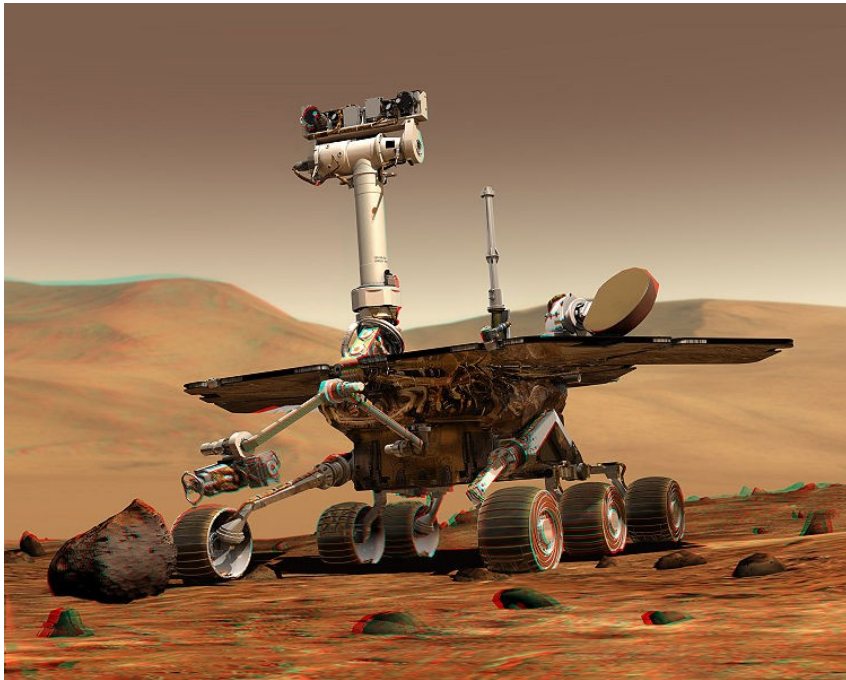
# Our time

- It is a good time to do computer vision now, because:
  - Powerful computers
  - Inexpensive cameras
  - Algorithm improvements
  - Understanding of vision systems

# Computer vision's applications

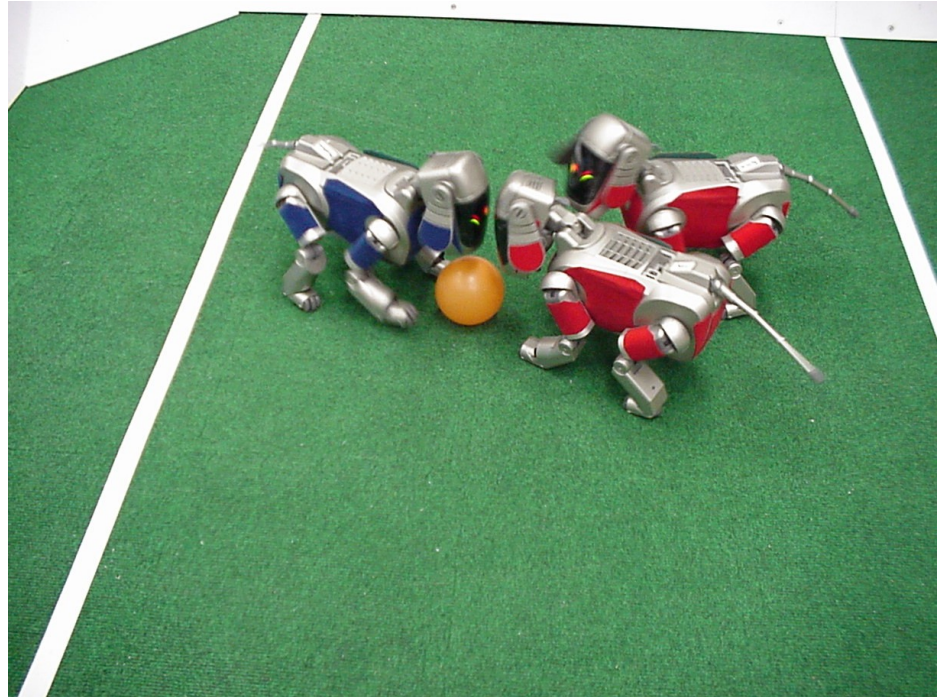
# Robotics

---



NASA's Mars Spirit Rover

[http://en.wikipedia.org/wiki/Spirit\\_rover](http://en.wikipedia.org/wiki/Spirit_rover)



<http://www.robocup.org/>

Adapted from S. Seitz

# 3D Maps

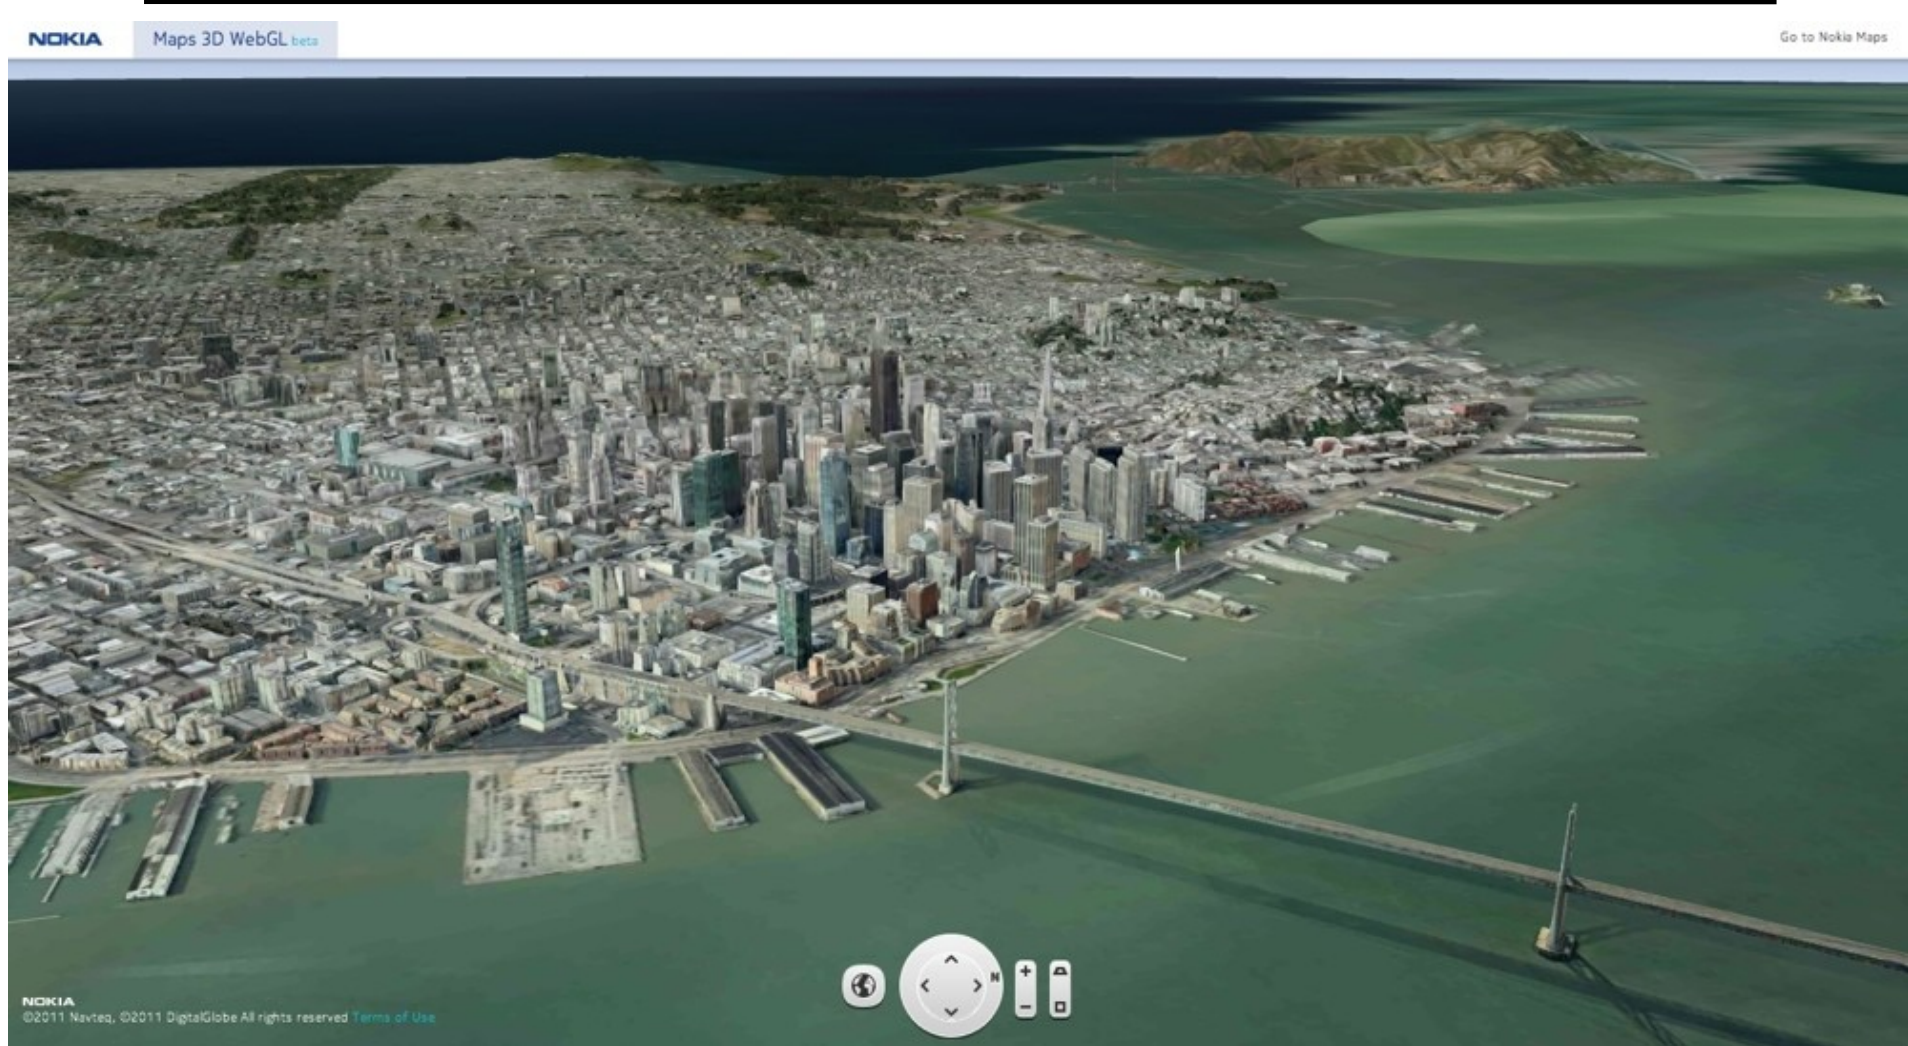


Image from Nokia's [Maps 3D WebGL](#)  
(see also: [Google Maps GL](#), [Google Earth](#))

Adapted from S. Seitz



# Motion capture

---



Microsoft's XBox Kinect

Adapted from S. Seitz

# Face detection

---



Most digital cameras detect faces

# Object recognition



Google Goggles  
Bing Vision

# Special effects: shape capture

---



The Matrix movies, ESC Entertainment, XYZRGB, NRC

Adapted from S. Seitz



# Sports

---



Sportvision first down line  
Nice [explanation](http://www.howstuffworks.com) on [www.howstuffworks.com](http://www.howstuffworks.com)

Adapted from S. Seitz

# Smart cars

Slide content courtesy of Amnon Shashua

The screenshot displays the Mobileye website layout. At the top, there are navigation tabs for 'manufacturer products' and 'consumer products'. The main header reads 'Our Vision. Your Safety.' Below this is a top-down view of a car with four camera fields of view highlighted: 'rear looking camera', 'forward looking camera', and 'side looking camera'. To the right, a 'News' sidebar lists articles such as 'Mobileye Advanced Technologies Power Volvo Cars World First Collision Warning With Auto Brake System' and 'Volvo: New Collision Warning with Auto Brake Helps Prevent Rear-end'. Below the main header, there are three product/application tiles: 'EyeQ Vision on a Chip' with an image of the chip, 'Vision Applications' showing a pedestrian detection box, and 'AWS Advance Warning System' with a dashboard display showing a car icon and a distance of 0.8. Each tile has a '> read more' link. A fourth 'Events' sidebar on the right lists 'Mobileye at Equip Auto, Paris, France' and 'Mobileye at SEMA, Las Vegas, NV', also with a '> read more' link.

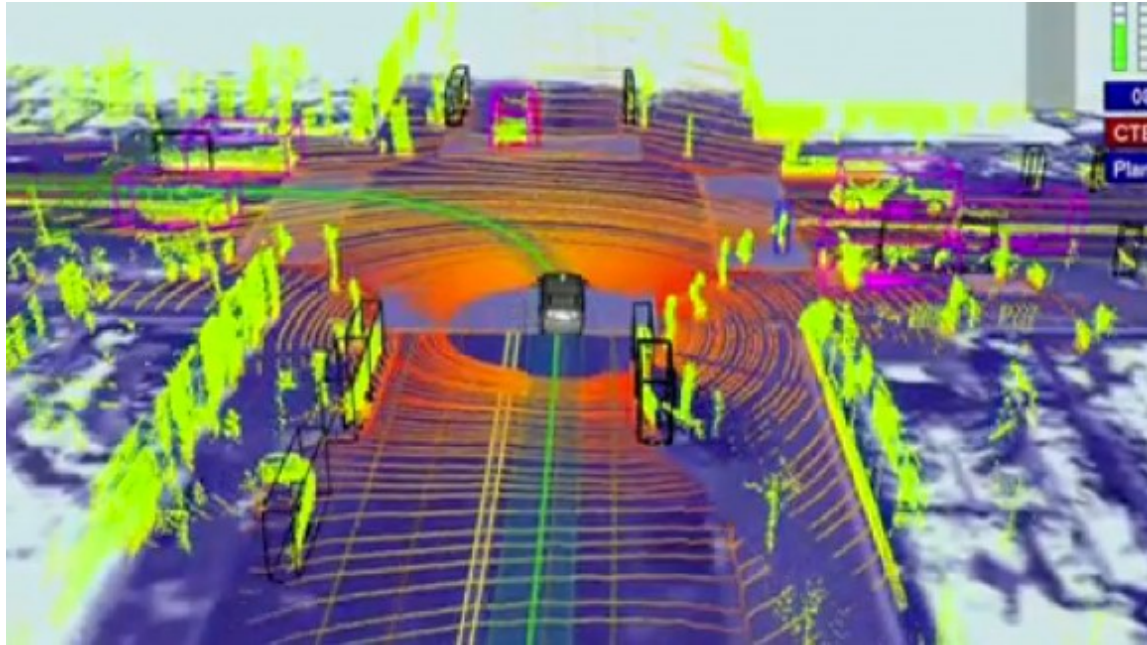
## Mobileye

- Vision systems currently in high-end BMW, GM, Volvo models

Adapted from S. Seitz

# Self-driving cars

---



“Our self-driving cars have now traveled nearly 200,000 miles on public highways in California and Nevada, 100 percent safely. They have driven from San Francisco to Los Angeles and around Lake Tahoe, and have even descended crooked Lombard Street in San Francisco. They drive anywhere a car can legally drive.”

- Sebastian Thrun, Google

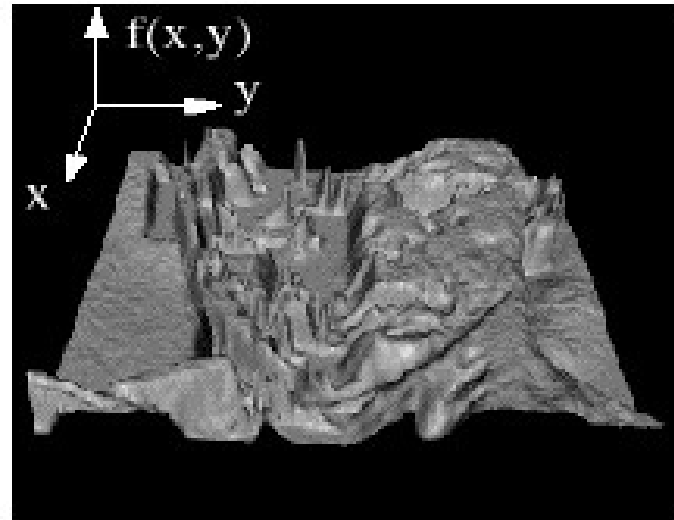
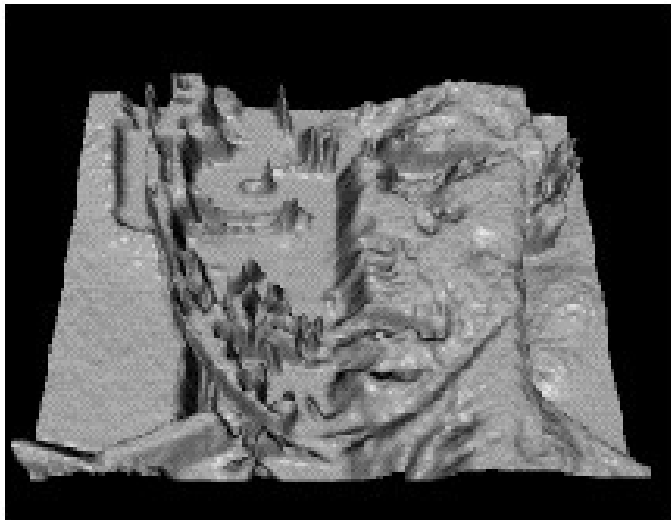
Adapted from S. Seitz

# Computer vision's applications

- To learn more about vision applications and companies
  - David lowe maintains an excellent overview of vision companies
    - <http://www.cs.ubc.ca/spider/lowe/vision.html>

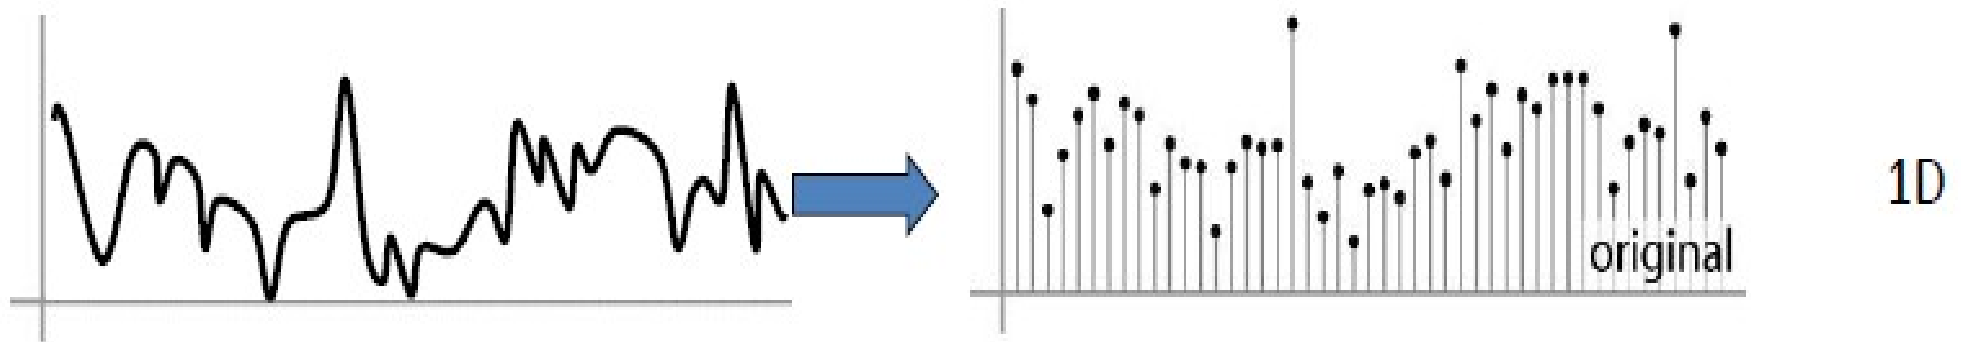
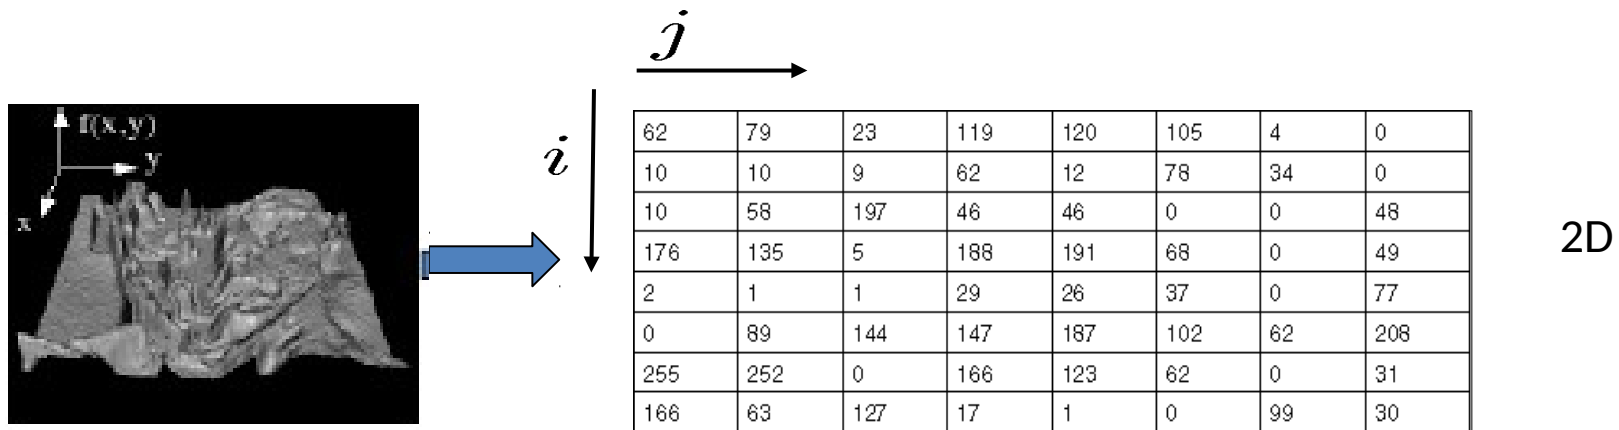
In the beginning ...

# Images as functions



# Digital images

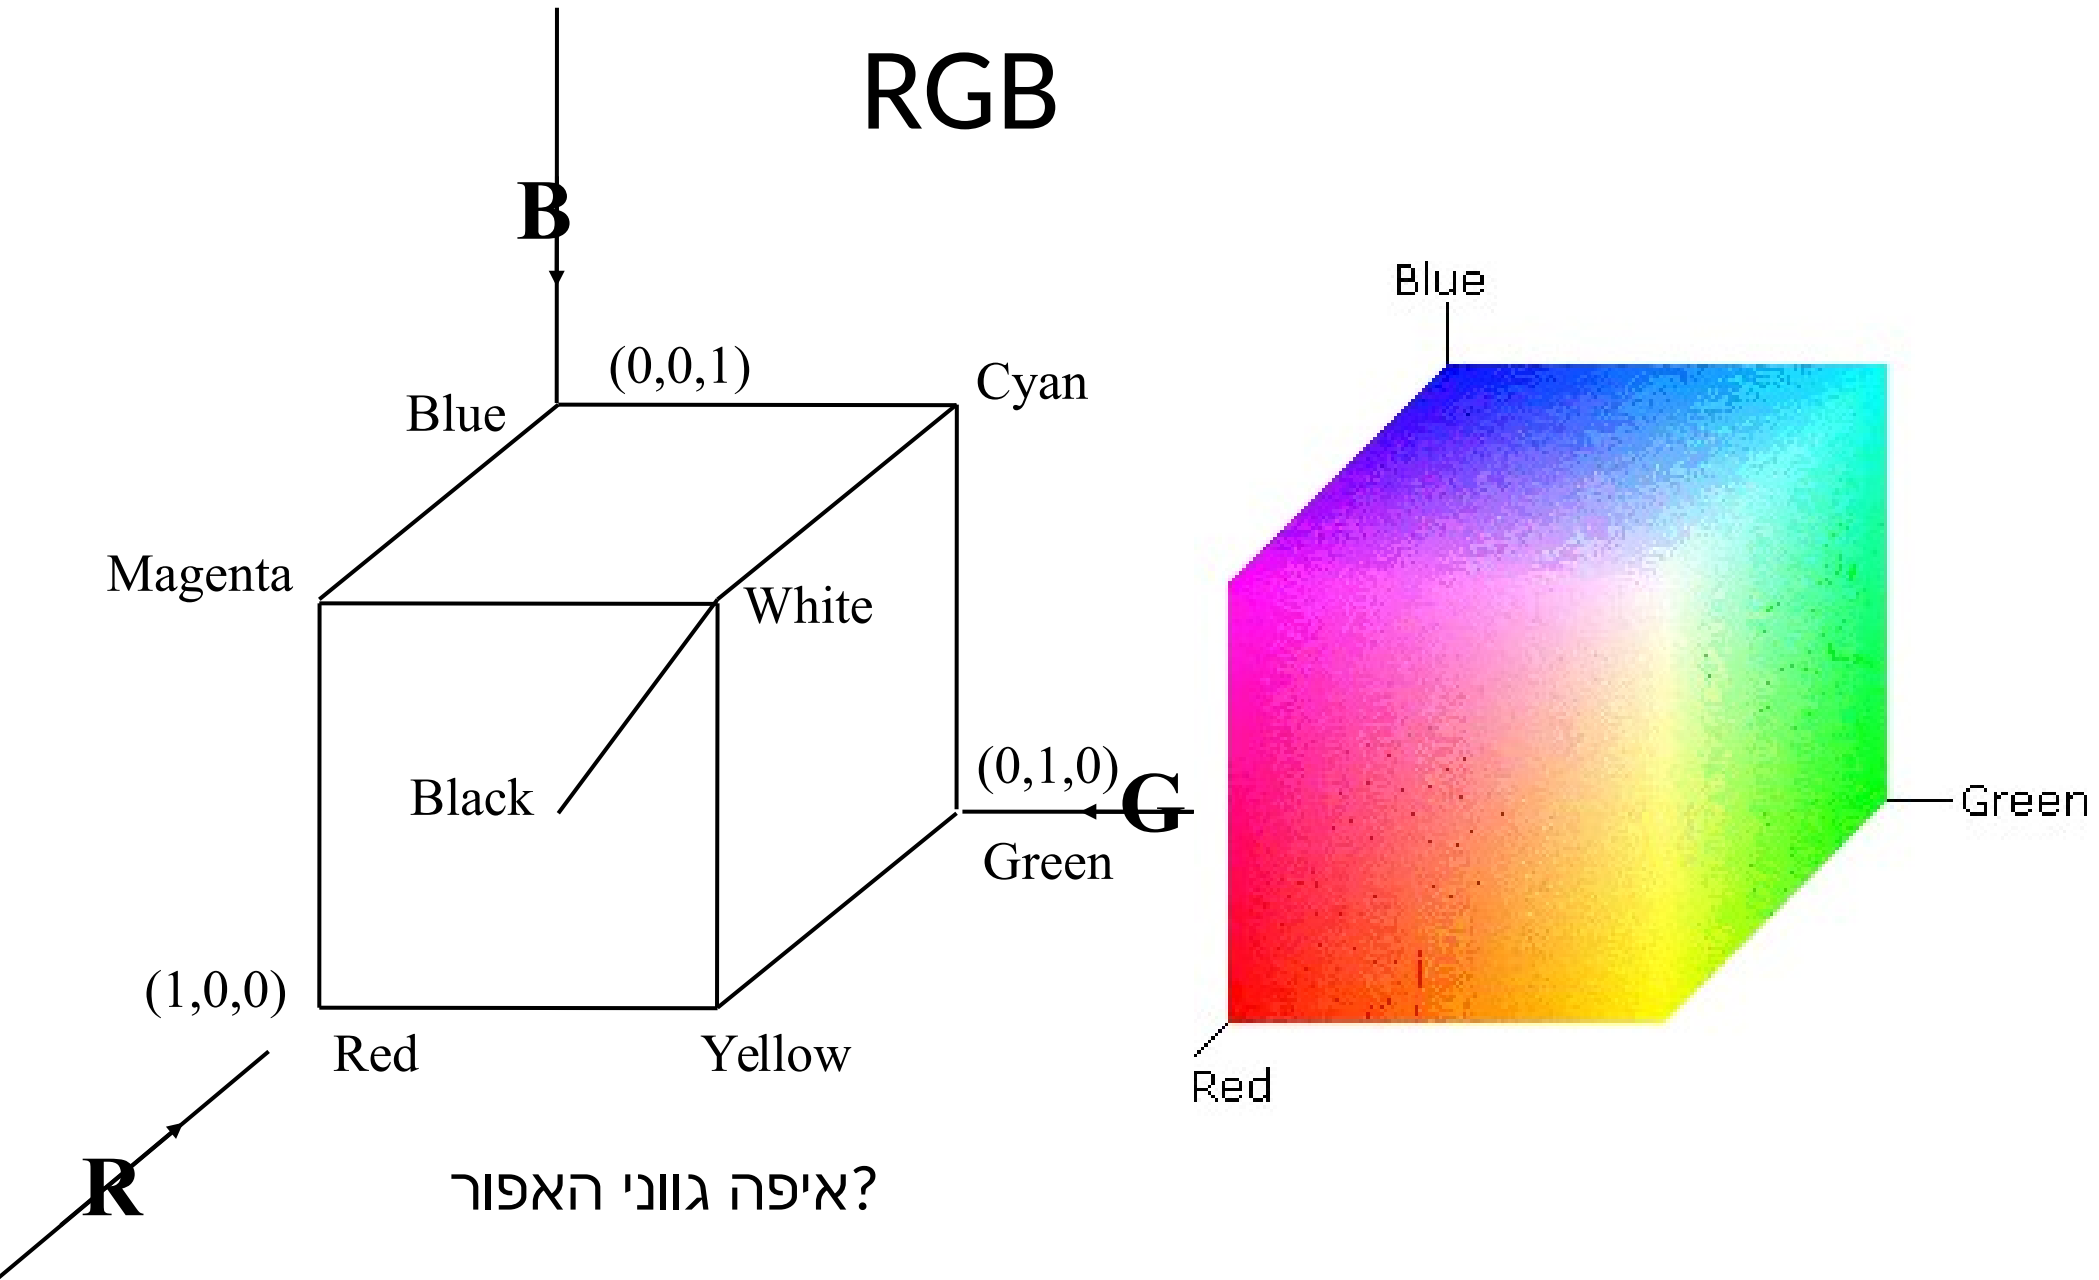
- In computer vision we operate on **digital (discrete)** images:
  - **Sample** the 2D space on a regular grid
  - **Quantize** each sample (round to nearest integer)
- Image thus represented as a matrix of integer values.



צבע



# RGB



# Color Spaces

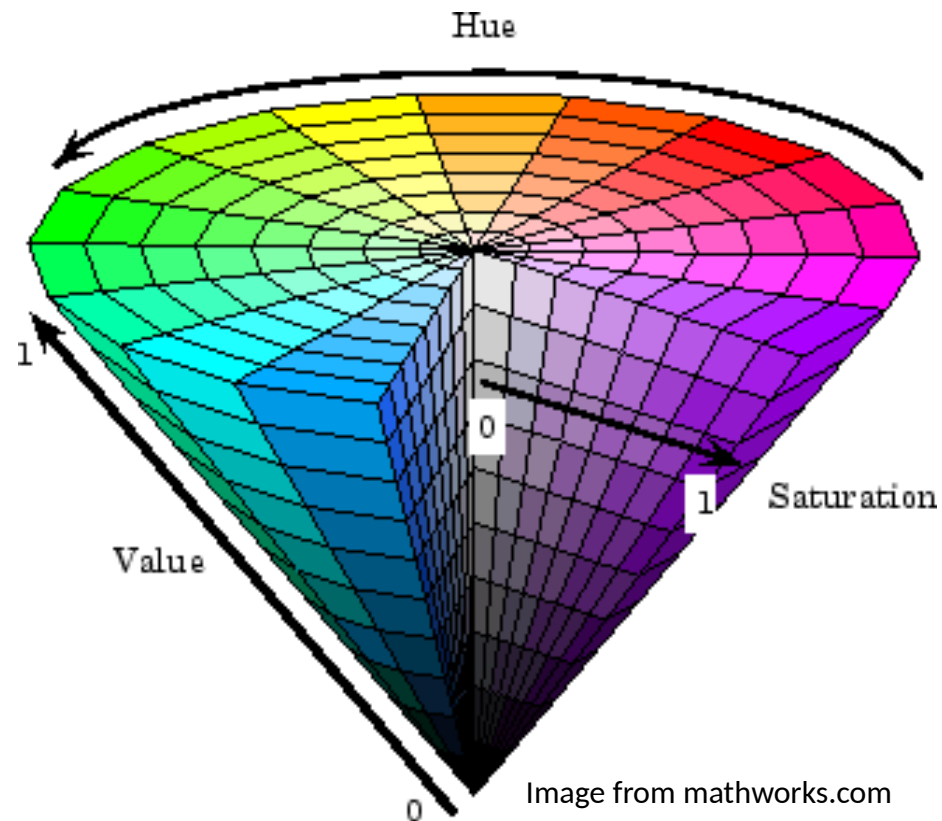
- “facilitate the specification of colors in some standard, generally accepted way. In essence, [...] a specification of a coordinate system and a subspace within that system where each color is represented by a single point”

Gonzalez & Woods

- RGB, CIE-XYZ, HSV, YIQ, CIE-Lab, YCbCr, ...

# HSV color space

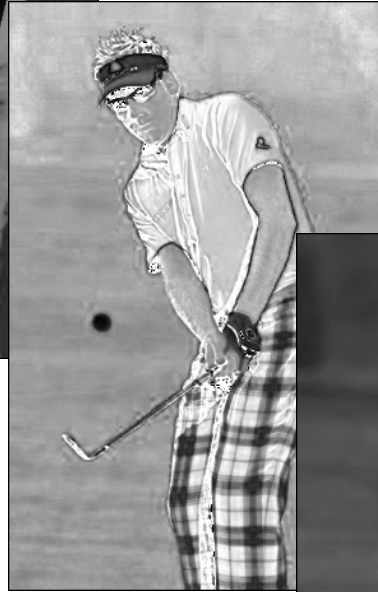
- **H**ue, **S**aturation, **V**alue (Brightness)
- Nonlinear – reflects topology of colors by coding **hue** as an angle
- Matlab: `hsv2rgb`,
- `rgb2hsv`.



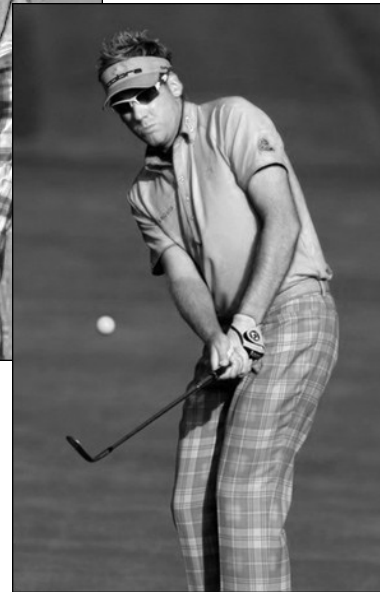
# HSV



H



S



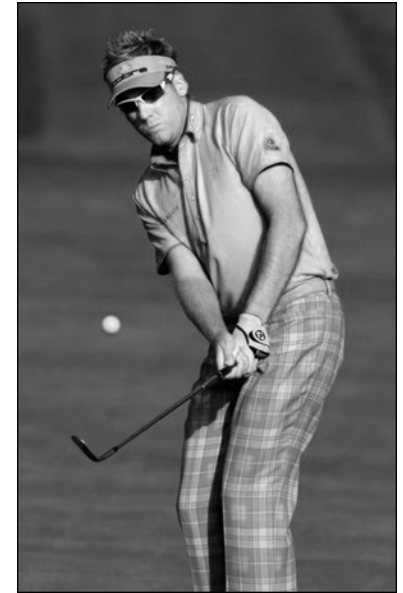
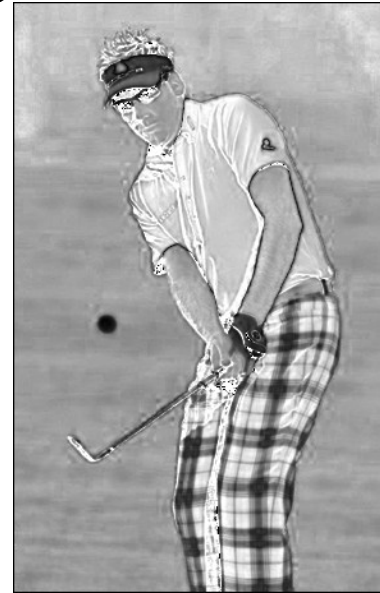
V

# HSV

H

S

V



## Python:

```
img = cv2.imread("golf.png")  
# Convert BGR to HSV  
hsv = cv2.cvtColor(img, cv2.COLOR_BGR2HSV)  
plt.imshow(hsv[:, :, 0], cmap='hsv')
```

**רעש ופילטרים**

# Common types of noise

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise



Impulse noise



Gaussian noise

# Linear filters

- Form a new image whose pixels are a weighted sum of the original pixel values, using the same set of weights at each point



# Box filter

- Mask with positive entries, that sum to 1.
- Replace each pixel with an average of its neighborhood.
- If all weights are equal, it is called a **box filter**.

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

# Correlation filtering

Say the averaging window size is  $(2k+1) \times (2k+1)$ :

$$G[i, j] = \underbrace{\frac{1}{(2k+1)^2}}_{\substack{\text{Attribute uniform} \\ \text{weight to each} \\ \text{pixel}}} \underbrace{\sum_{u=-k}^k \sum_{v=-k}^k F[i+u, j+v]}_{\substack{\text{Loop over all pixels in neighborhood} \\ \text{around image pixel } F[i,j]}}$$

Now generalize to allow different weights depending on neighboring pixel's relative position:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k \underbrace{H[u, v]}_{\text{Non-uniform weights}} F[i+u, j+v]$$

# Correlation filtering

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

This is called cross-correlation, denoted

$$G = H \otimes F$$

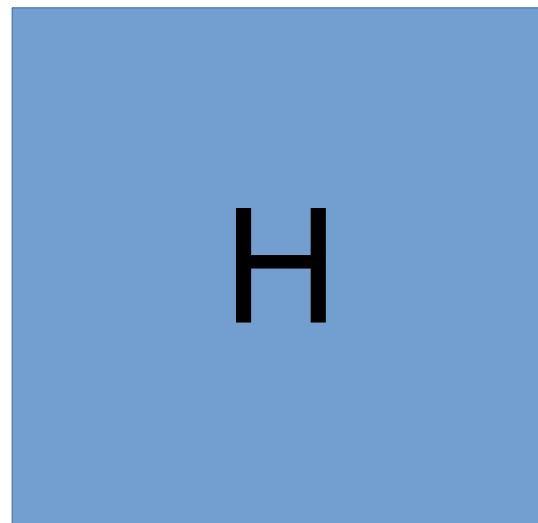
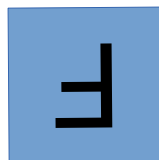
Filtering an image: replace each pixel with a linear combination of its neighbors.

*The filter “kernel” or “mask”  $H[u,v]$  is the prescription for the weights in the linear combination.*

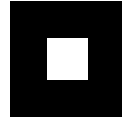
# Convolution

$$\begin{aligned} G[i, j] &= H * F \\ &= \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v] \end{aligned}$$

- Flip the filter in both dimension.



# Smoothing by averaging



depicts box filter:  
white = high value, black = low value

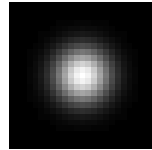


original



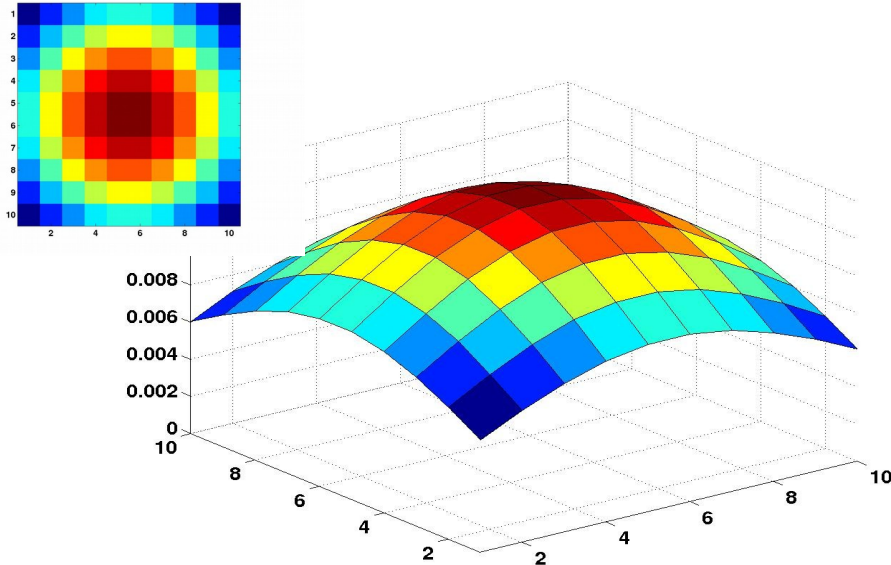
filtered

# Smoothing with a Gaussian

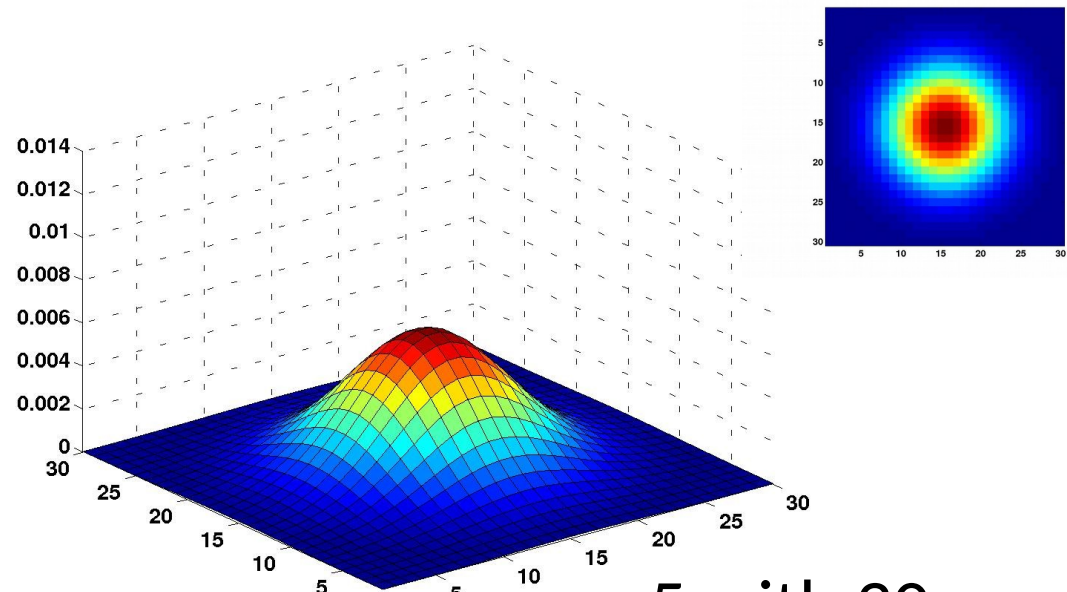


# Gaussian filters

- What parameters matter here?
- **Size** of kernel or mask
  - Note, Gaussian function has infinite support, but discrete filters use finite kernels



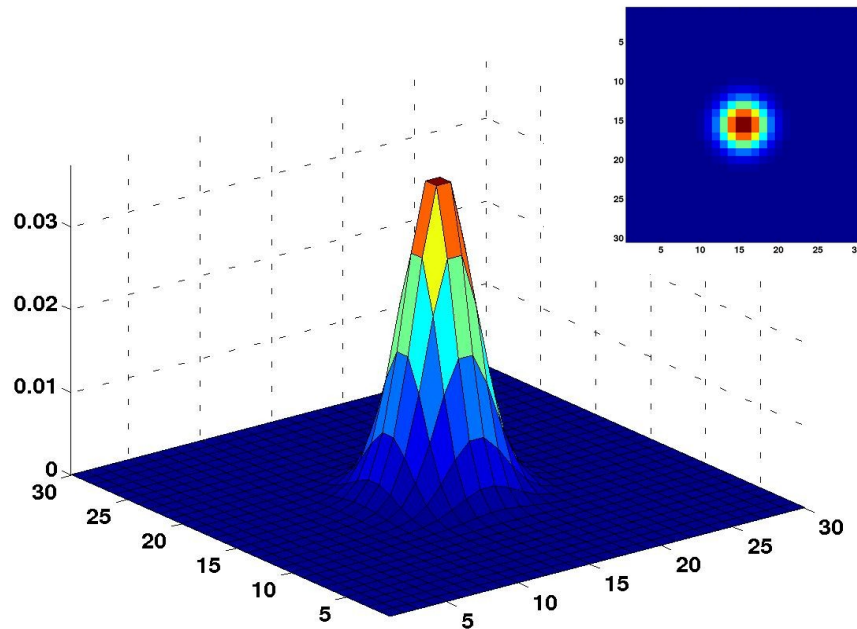
$\sigma = 5$  with 10  
x 10 kernel



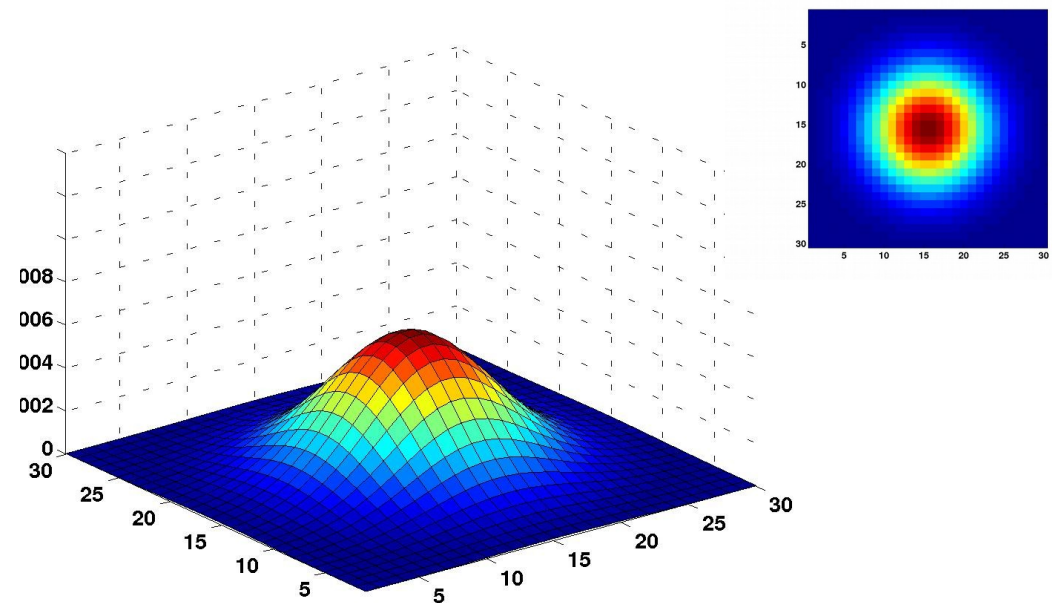
$\sigma = 5$  with 30  
x 30 kernel

# Gaussian filters

- What parameters matter here?
- **Variance** of Gaussian: determines extent of smoothing



$\sigma = 2$  with 30  
x 30 kernel



$\sigma = 5$  with 30  
x 30 kernel



# Efficient Implementation - Separability

- In some cases, filter is separable,

$$\begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array}$$

- Thus, using the associativity property,

$$(f * g) * h = f * (g * h)$$

- compute in two steps.

# Separability of the Gaussian filter

- The 2D Gaussian is separable:

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2+y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi\sigma^2}} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi\sigma^2}} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

- What about the Box-filter?

# Properties of convolution

- Linear & shift invariant

- Commutative:  $f * g = g * f$

- Associative

- Identity:  $(f * g) * h = f * (g * h)$

- unit impulse

- Differentiation:  $e = [\dots, 0, 0, 1, 0, 0, \dots] \Rightarrow f * e = f$

$$\frac{\partial}{\partial x}(f * g) = \frac{\partial f}{\partial x} * g$$

# Effect of smoothing filters

Gaussian  
noise

Salt and pepper  
noise

3x3



5x5

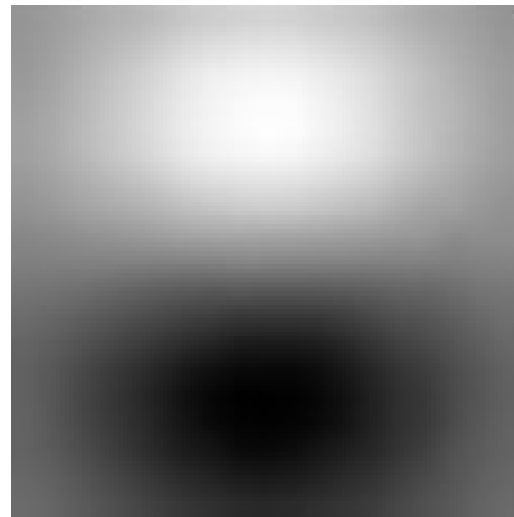
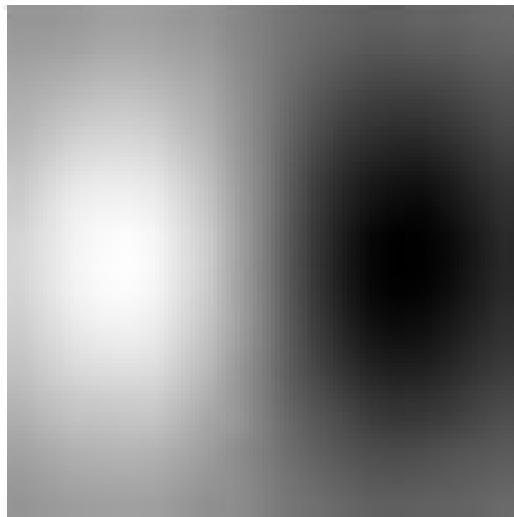


7x7

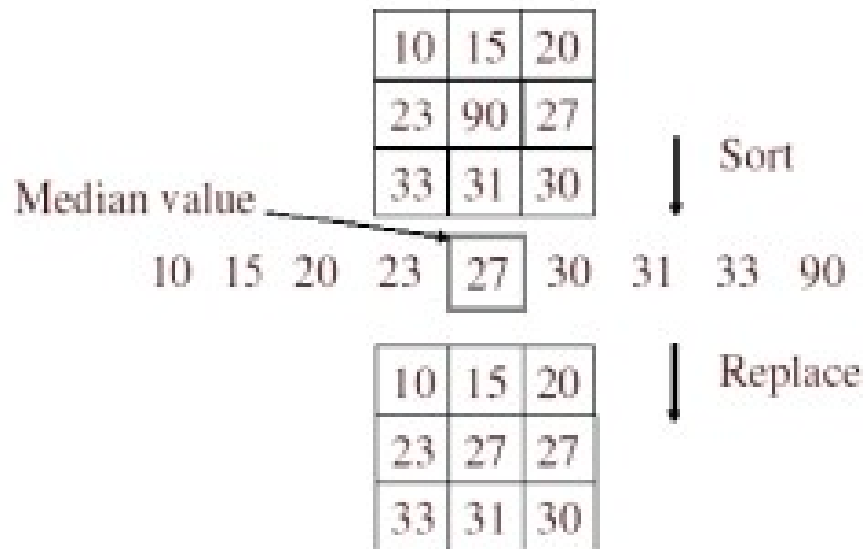


# Filter as templates

- Applying a filter can be seen as taking dot-product between an image area and a vector.
- Filter emphasizes similar areas.
  - Highest response for regions that “look the most like the filter”



# Median filter



- Nonlinear
  - No new pixel values introduced
  - Remove spikes: good for impulse, salt&pepper noise
  - Drawback: produces unexpected artifacts, erase fine lines and/or inverts parallel fine lines black → white and white → black, erodes sharp edges, slow.

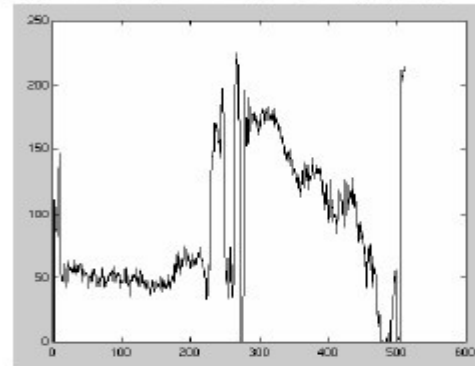
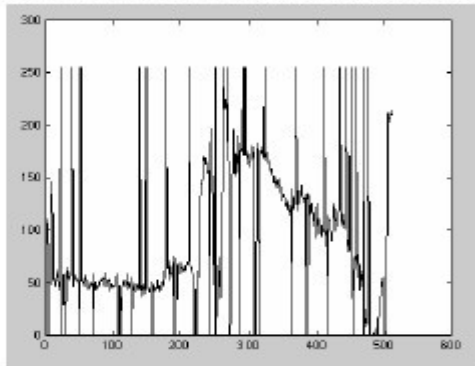
# Median filter

Effect of median filter on salt and pepper noise

Salt and  
pepper noise →



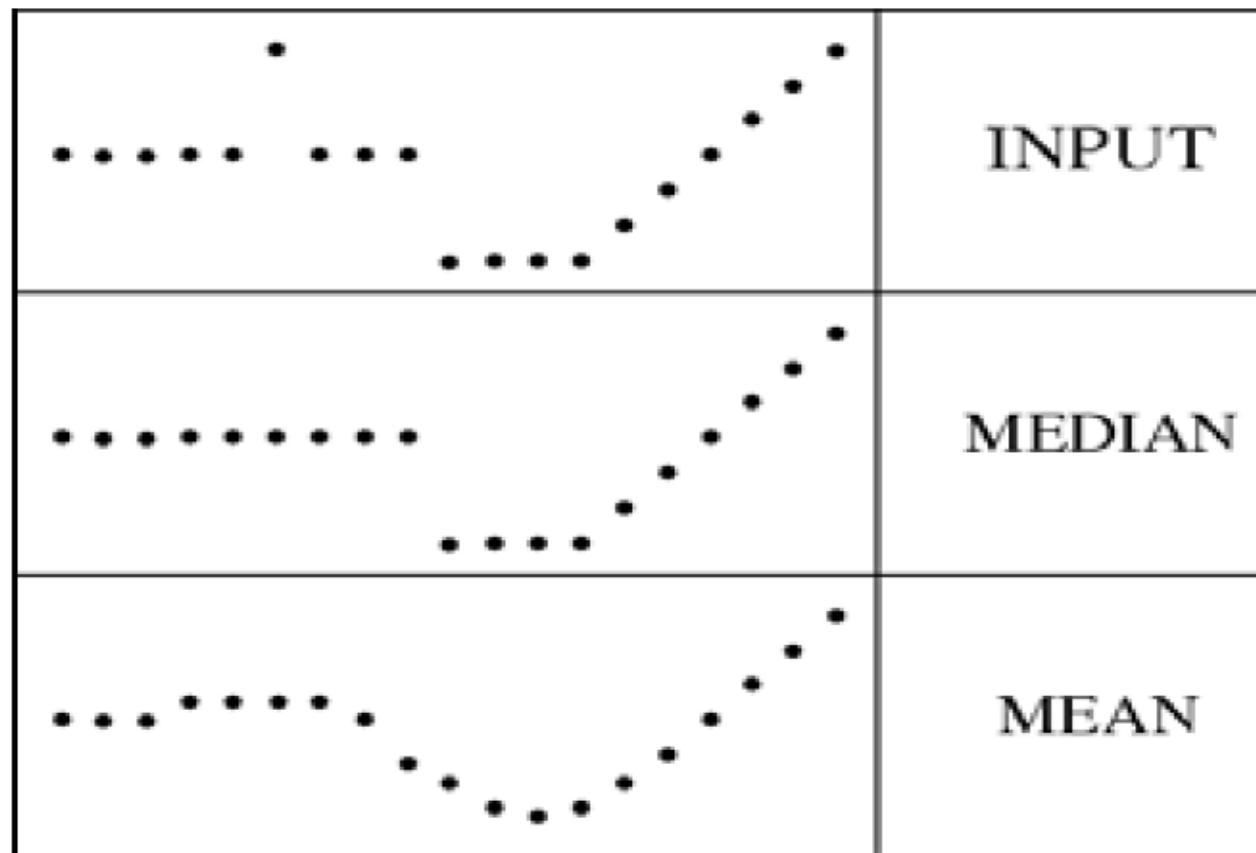
← Median  
filtered



Plots of a row of the image

# Median filter

- Median filter is edge preserving





# Edge detection

- **Goal:** map image from 2d array of pixels to a set of curves or line segments or contours.
- **Why?**

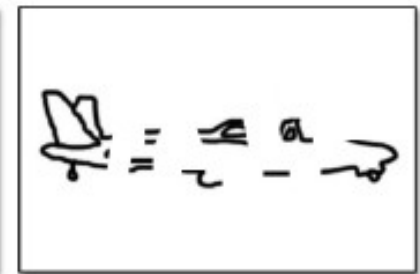
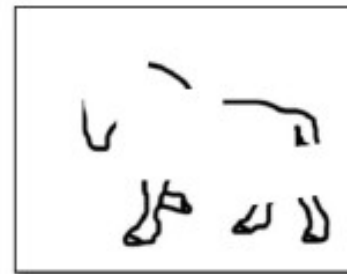
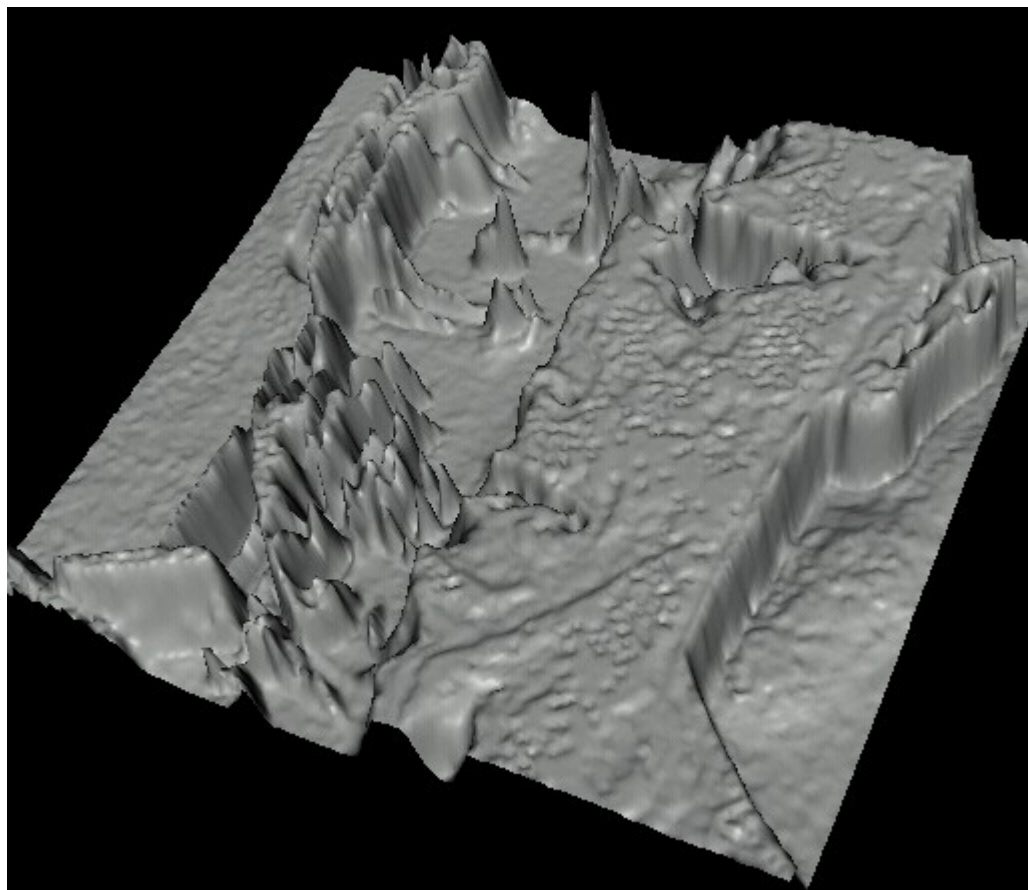


Figure from J. Shotton et al., PAMI 2007

- **Main idea:** look for strong gradients, post-process

# Recall : Images as functions



- Edges look like steep cliffs

Source: S. Seitz

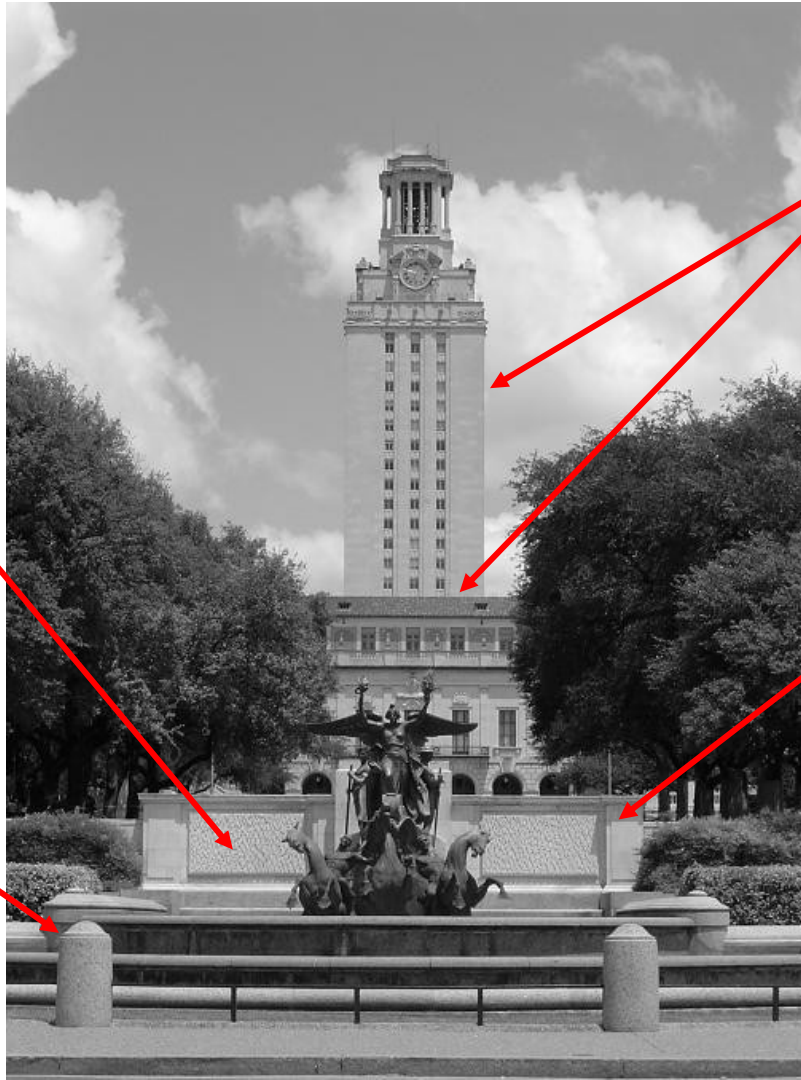
# What can cause an edge?

Reflectance change:  
appearance  
information, texture

Depth discontinuity:  
object boundary

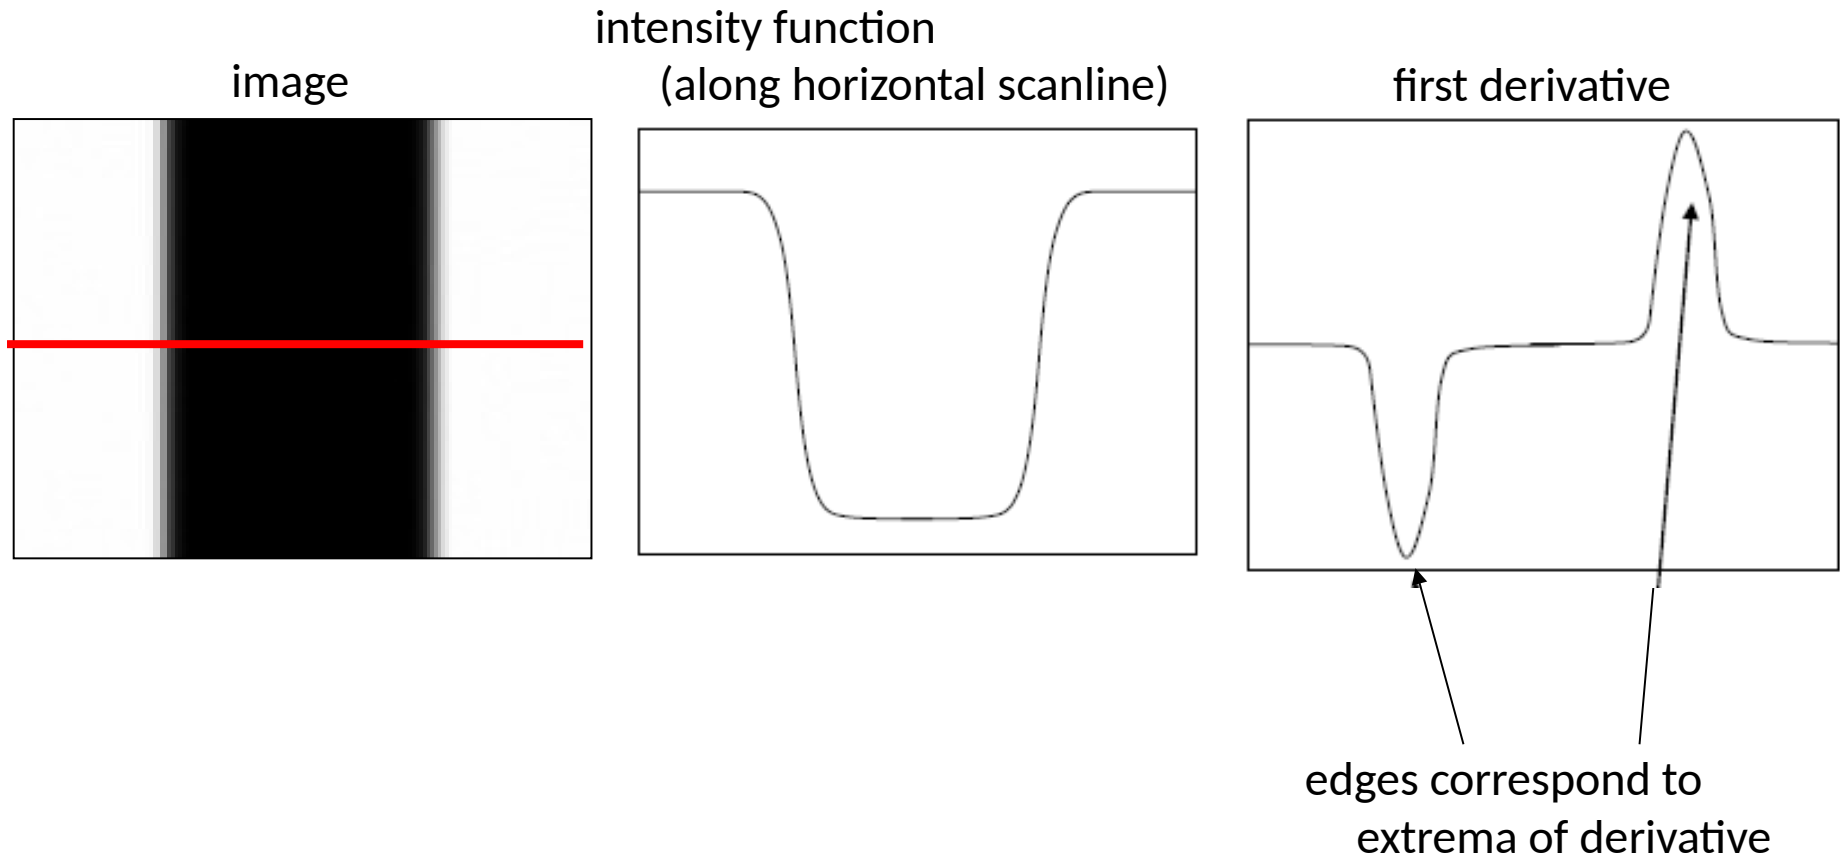
Change in surface  
orientation: shape

Cast shadows



# Derivatives and edges

An edge is a place of rapid change in the image intensity function.



# Differentiation and convolution

For 2D function,  $f(x,y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon, y) - f(x, y)}{\epsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

To implement above as convolution, what would be the associated filter?

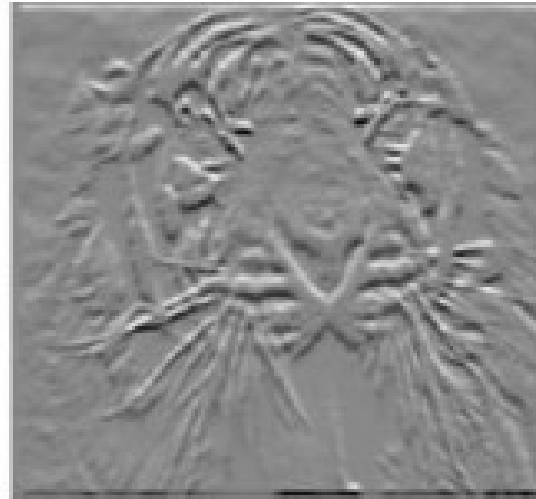
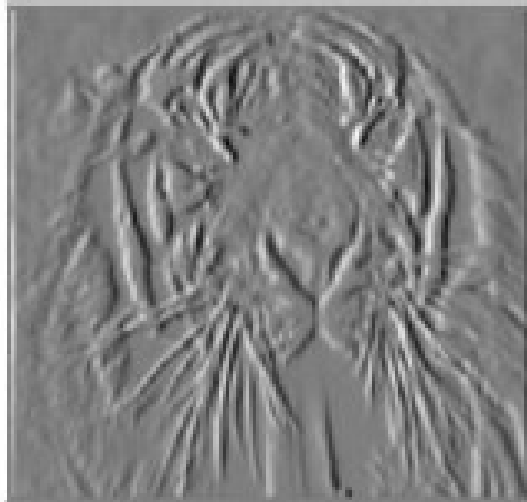
# Partial derivatives of an image



$$\frac{\partial f(x, y)}{\partial x}$$

$$\frac{\partial f(x, y)}{\partial y}$$

-1	1
----	---



-1
1

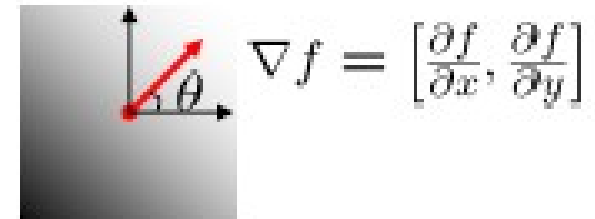
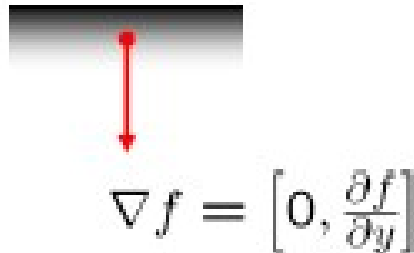
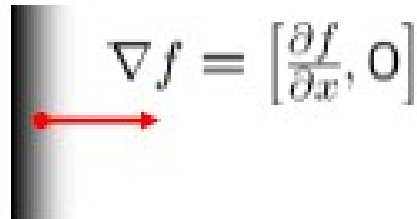
Which shows changes with respect to x?

# Image gradient

- The gradient of an image

$$\nabla f(x, y) = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

- The gradient points to the direction of the most rapid change in intensity



The gradient direction is the orientation of the edge normal:

$$\theta = \tan^{-1} \left( \frac{\frac{\partial f}{\partial y}}{\frac{\partial f}{\partial x}} \right)$$

- The edge strength is given by the gradient magnitude:

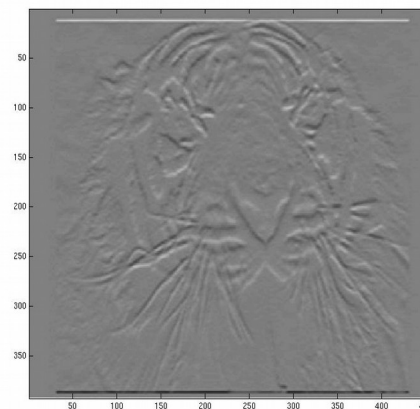
$$||\nabla f|| = \sqrt{\left( \frac{\partial f}{\partial x} \right)^2 + \left( \frac{\partial f}{\partial y} \right)^2}$$

# Assorted finite difference filters

**Prewitt:**  $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix} ; M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

**Sobel:**  $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} ; M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

**Roberts:**  $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix} ; M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

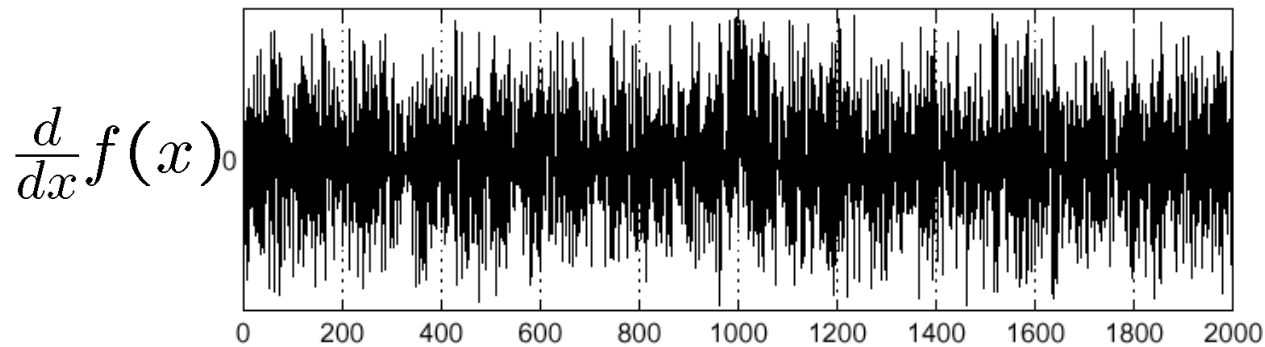
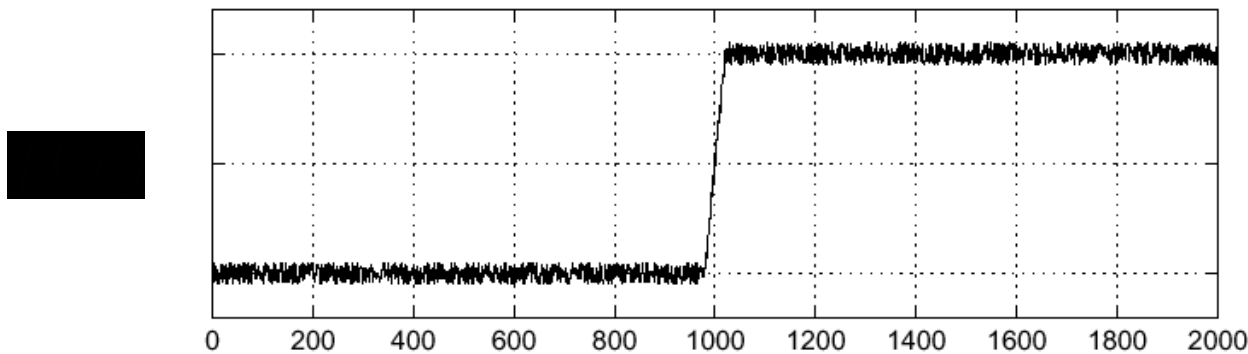




# Effects of noise

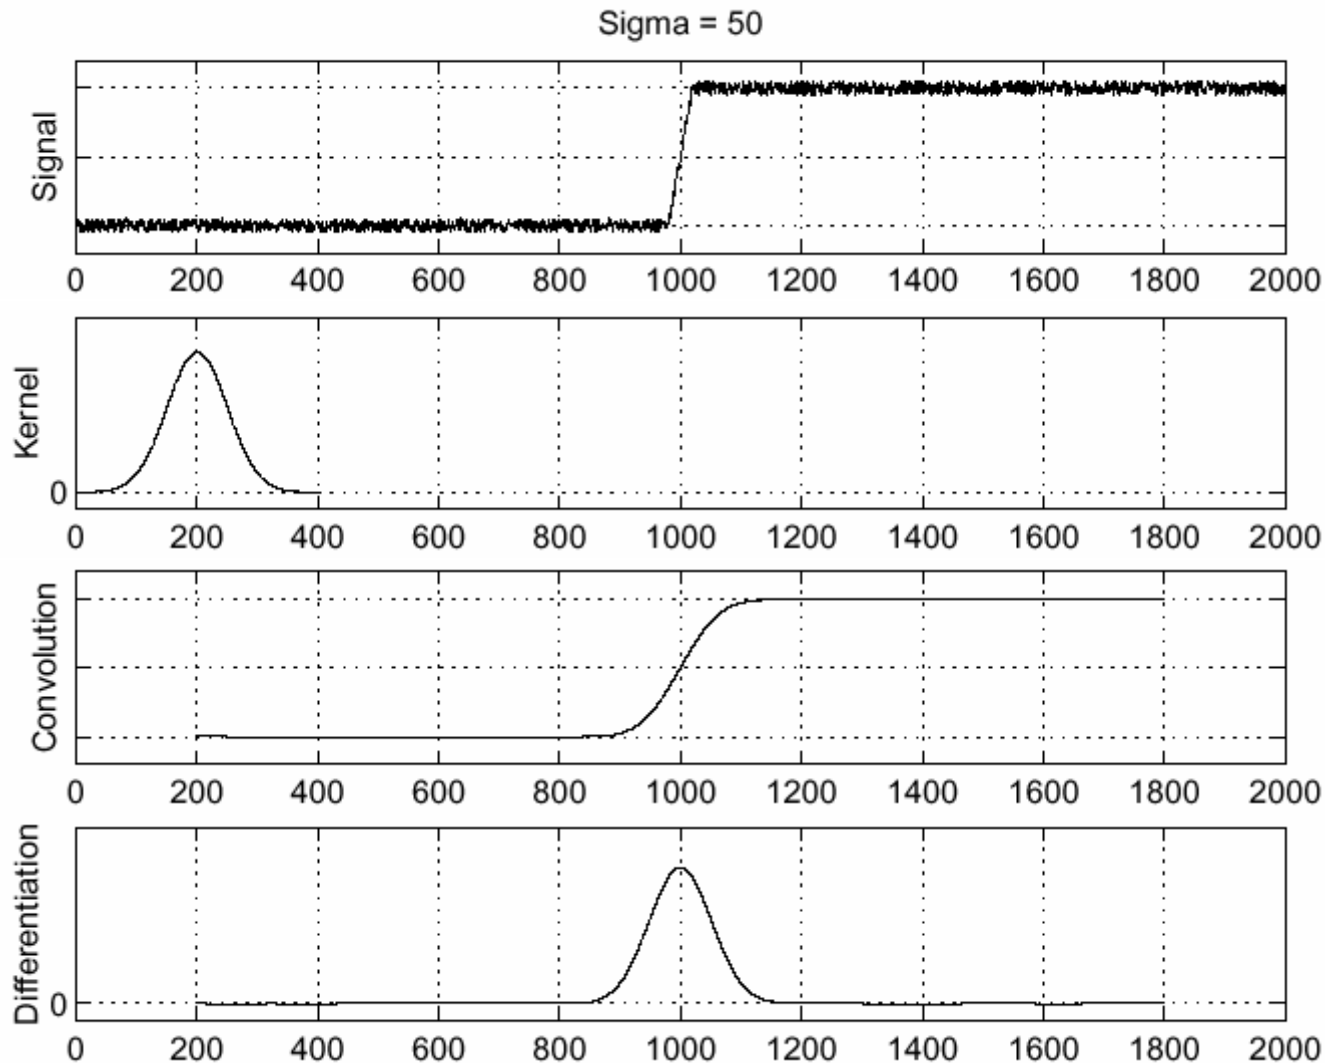
Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal



Where is the edge?

# Solution: smooth first



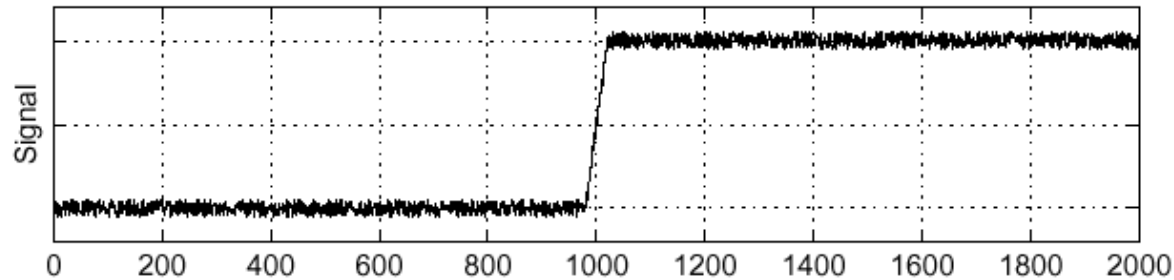
Where is the edge? Look for peaks in  $\frac{\partial}{\partial x}(h \star f)$

# Derivative theorem of convolution

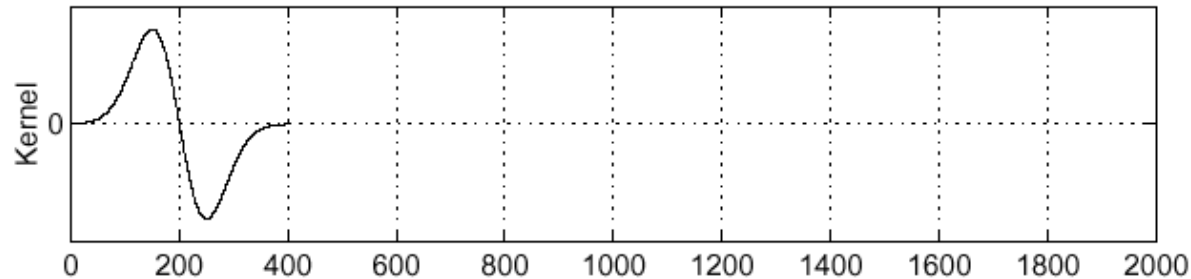
---

$$\frac{\partial}{\partial x}(h \star f) = \left(\frac{\partial}{\partial x}h\right) \star f$$

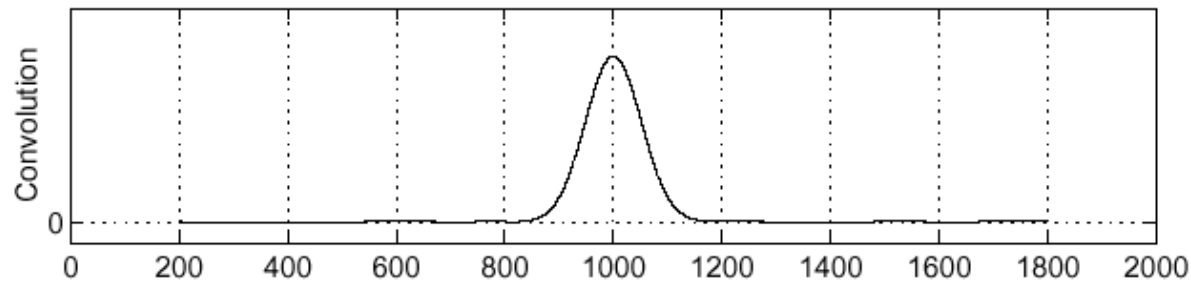
This saves us one operation:



$$\frac{\partial}{\partial x}h$$



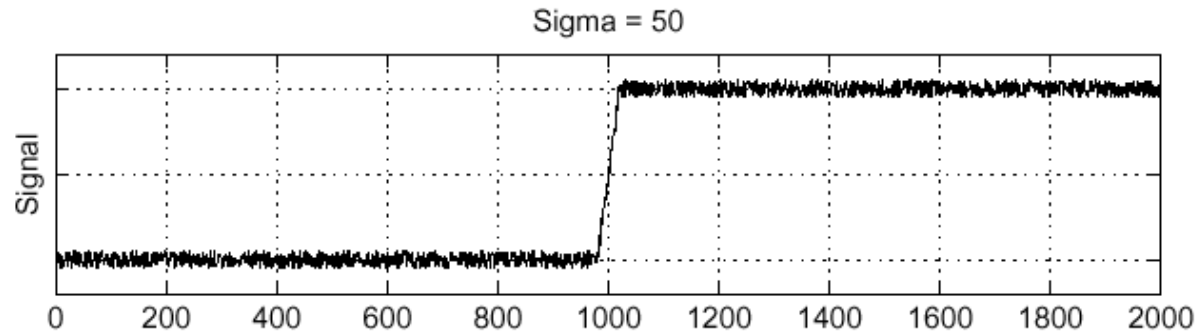
$$\left(\frac{\partial}{\partial x}h\right) \star f$$



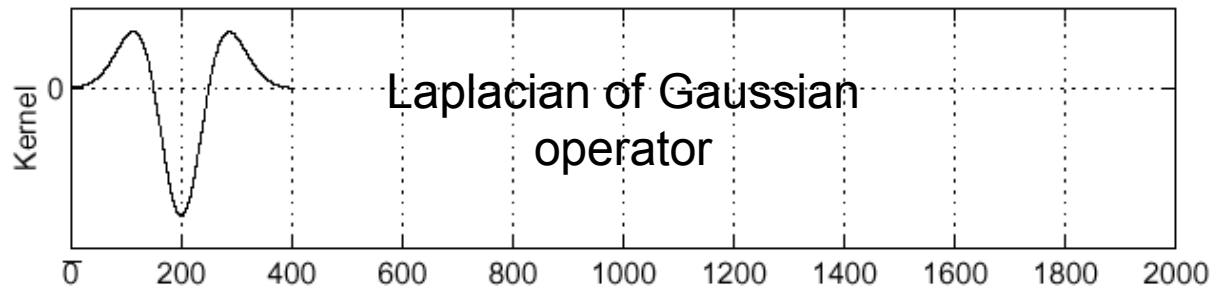
How can we find (local) maxima of a function?

# Laplacian of Gaussian

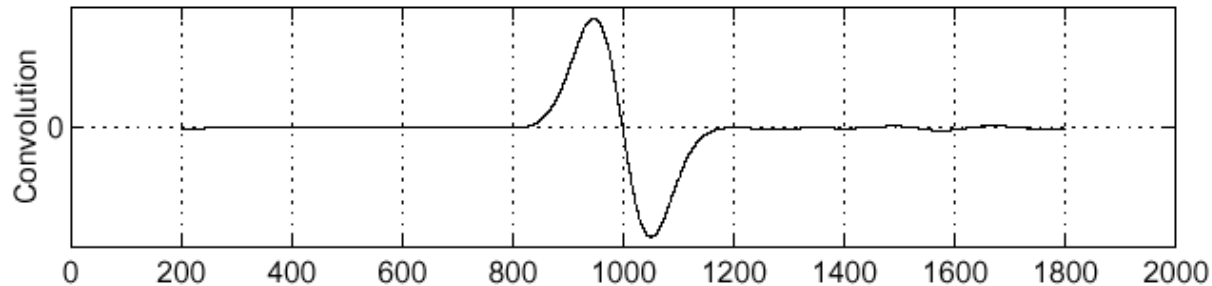
Consider  $\frac{\partial^2}{\partial x^2}(h \star f)$



$$\frac{\partial^2}{\partial x^2}h$$



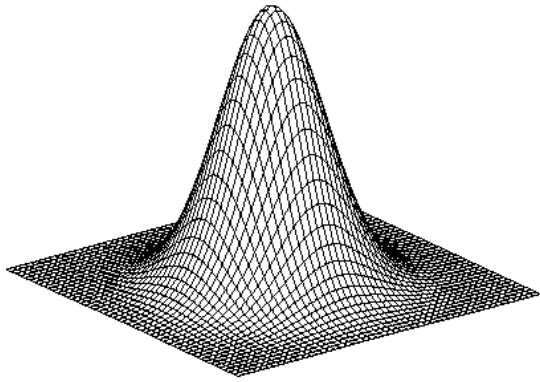
$$\left(\frac{\partial^2}{\partial x^2}h\right) \star f$$



Where is the edge?      Zero-crossings of bottom graph

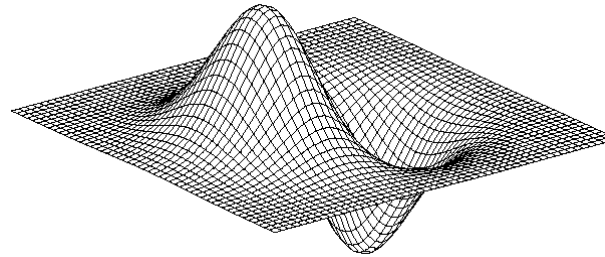
# 2D edge detection filters

---



Gaussian

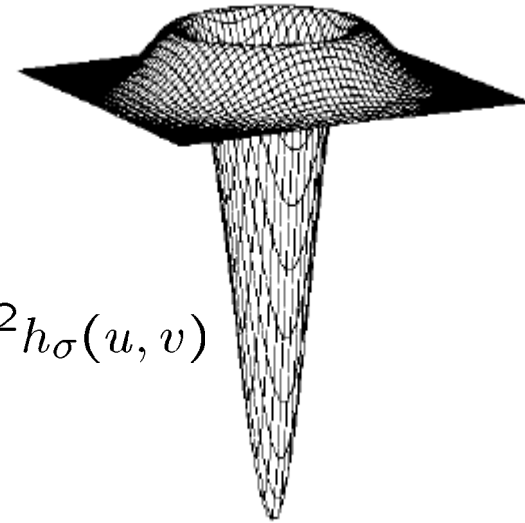
$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$

Laplacian of Gaussian

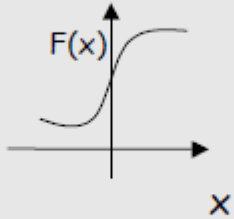
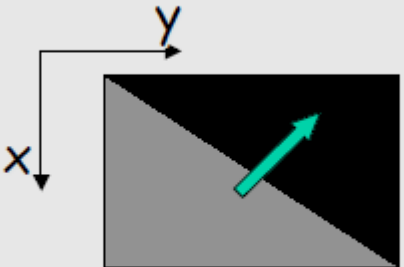


$$\nabla^2 h_{\sigma}(u, v)$$

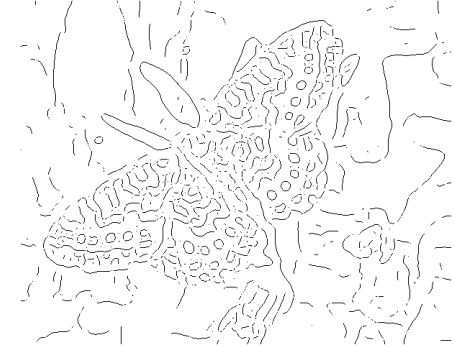
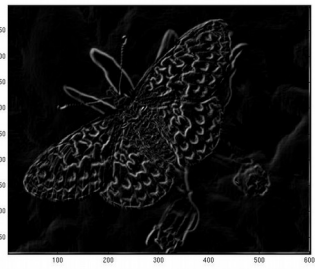
■ is the Laplacian operator:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

# Edge detection summary

	1D	2D
step edge	$I(x)$ 	$I(x,y)$ 
1st deriv	$\left  \frac{dI(x)}{dx} \right  > Th$	$ \nabla I(x,y)  = (I_x^2(x,y) + I_y^2(x,y))^{1/2} > Th$ $\tan \theta = I_x(x,y) / I_y(x,y)$
2nd deriv	$\frac{d^2I(x)}{dx^2} = 0$	$\nabla^2 I(x,y) = I_{xx}(x,y) + I_{yy}(x,y) = 0$ <div style="border: 1px solid black; border-radius: 50%; padding: 5px; display: inline-block;">Laplacian</div>

# Gradients -> edges



Primary edge detection steps:

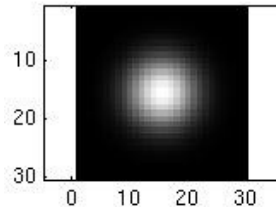
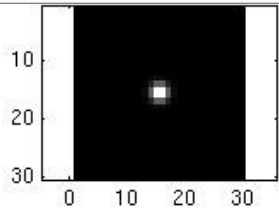
1. Smoothing: suppress noise
2. Edge enhancement: filter for contrast
3. Edge localization

Determine which local maxima from  
filter output are actually edges vs. noise

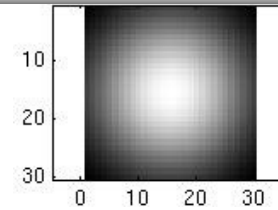
- Threshold, Thin

# Smoothing with a Gaussian

Recall: parameter  $\sigma$  is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



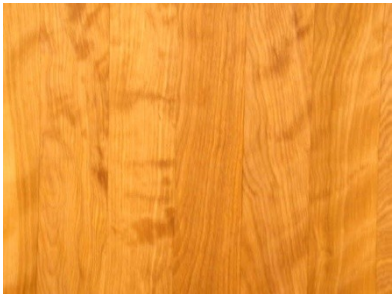
...





# So, what scale to choose?

It depends what we're looking for.



Too fine of a scale...can't see the forest for the trees.

Too coarse of a scale...can't tell the maple grain from the cherry.

# Canny edge detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- **Non-maximum suppression:**
  - Thin multi-pixel wide “ridges” down to single pixel width
- Linking and thresholding (**hysteresis**):
  - Define two thresholds: low and high
  - Use the high threshold to start edge curves and the low threshold to continue them

# The Canny edge detector



original image (Lena)

# The Canny edge detector



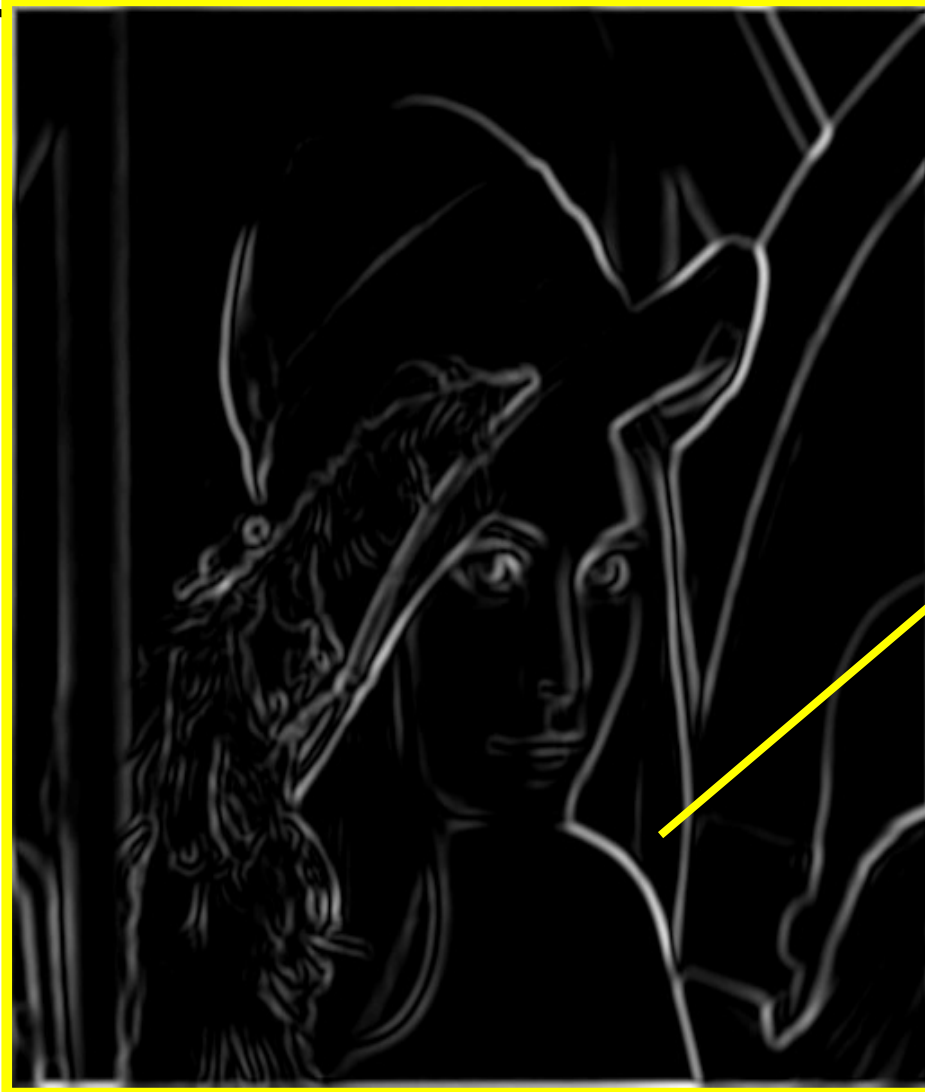
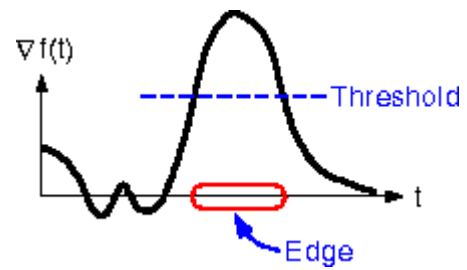
norm of the gradient

# The Canny edge detector



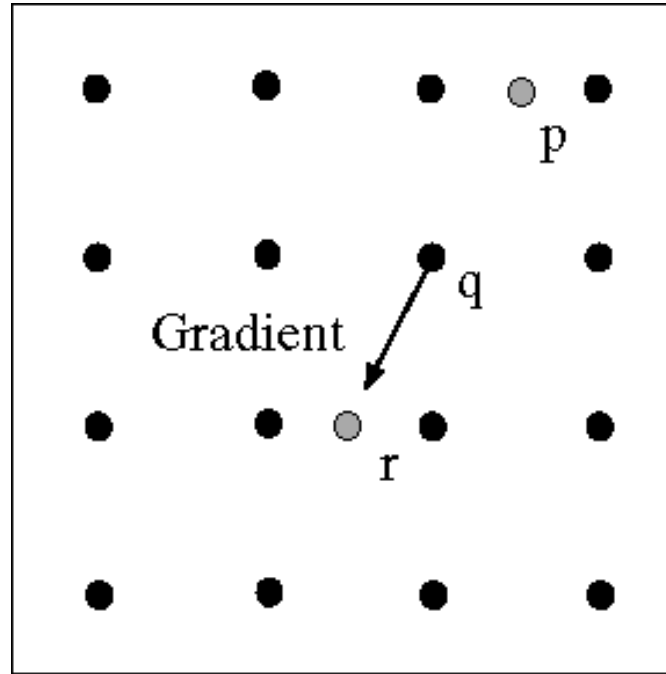
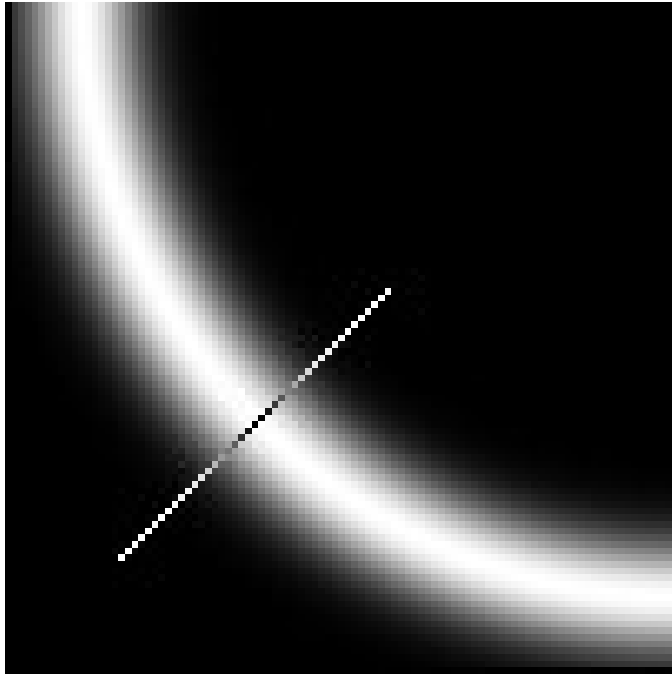
thresholding

# The Canny edge detector



How to turn these thick regions of the gradient into curves?

# Non-maximum suppression

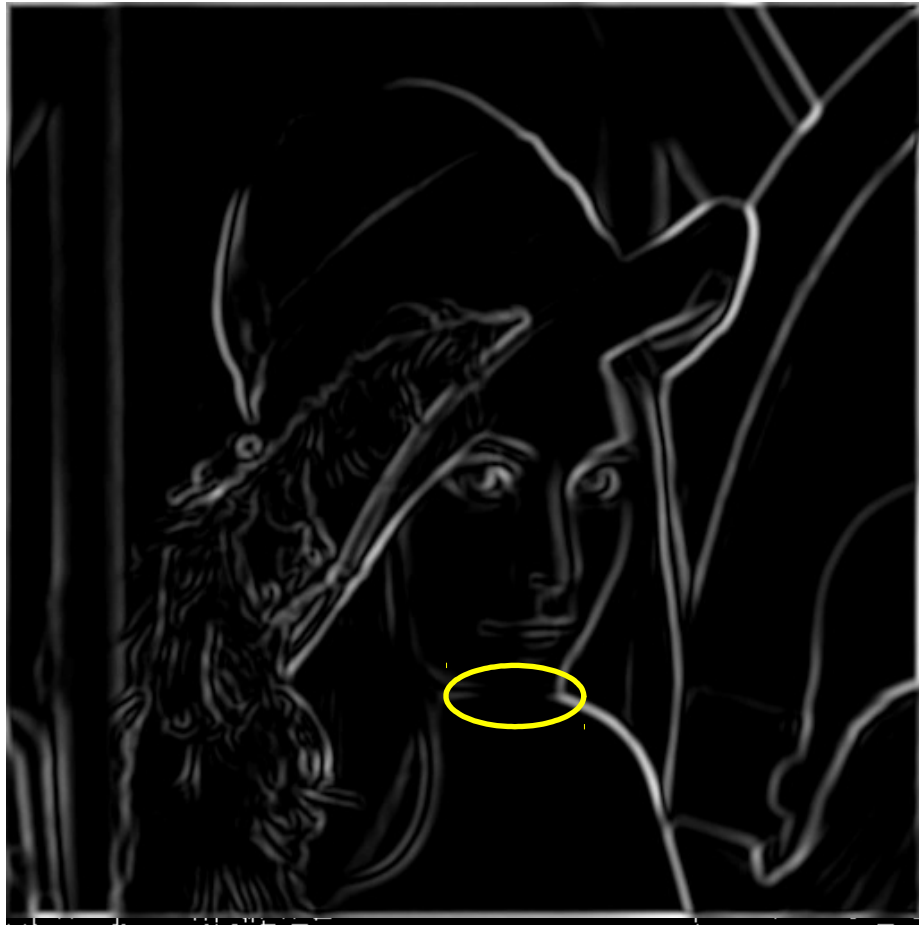


Check if pixel is local maximum along gradient direction,  
select single max across width of the edge

- requires checking interpolated pixels  $p$  and  $r$



# The Canny edge detector



Problem: pixels  
along this edge  
didn't survive  
the  
thresholding

thinning  
(non-maximum suppression)



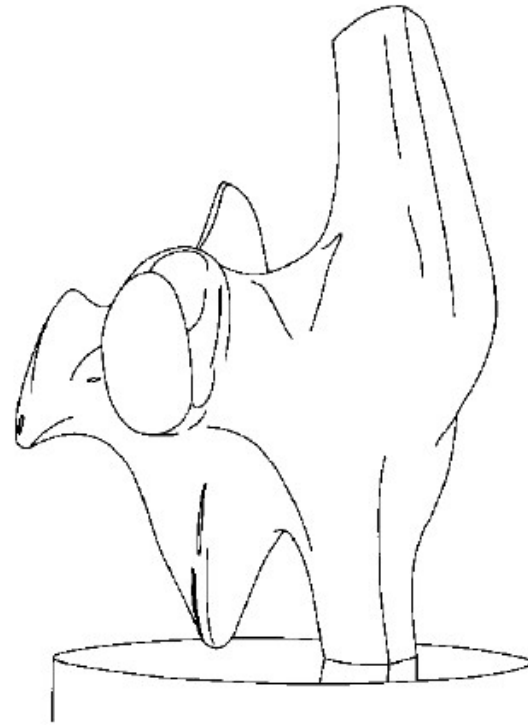
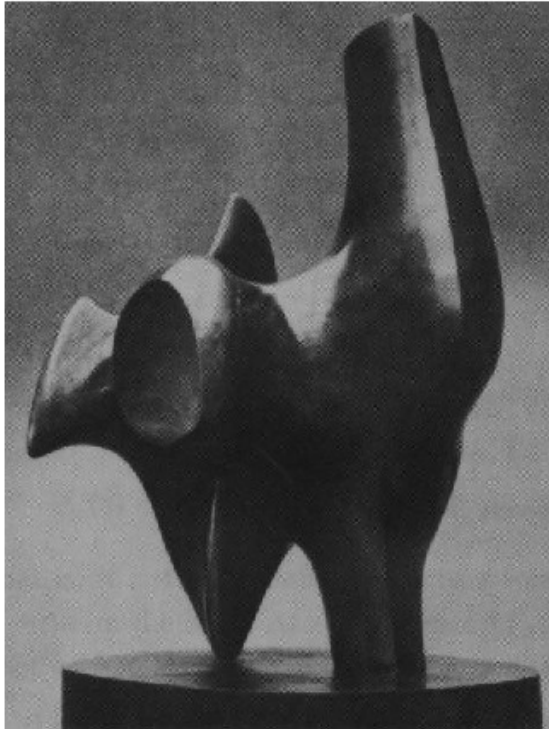
# Hysteresis thresholding

- Check that maximum value of gradient value is sufficiently large
- drop-outs? use **hysteresis**
  - use a high threshold to start edge curves and a low threshold to continue them.



Source: S. Seitz

# Summary



# Summary

