

# מבוא לראייה ממוחשבת – 22928 2016א

מנחה: אמיר אגוזי  
[egozi5@gmail.com](mailto:egozi5@gmail.com)

מפגש מס' 4

# בפעם שעברה

PCA •

עקרונות לימוד מכונה •

Feature extraction –

Training –

Validation –

Test –

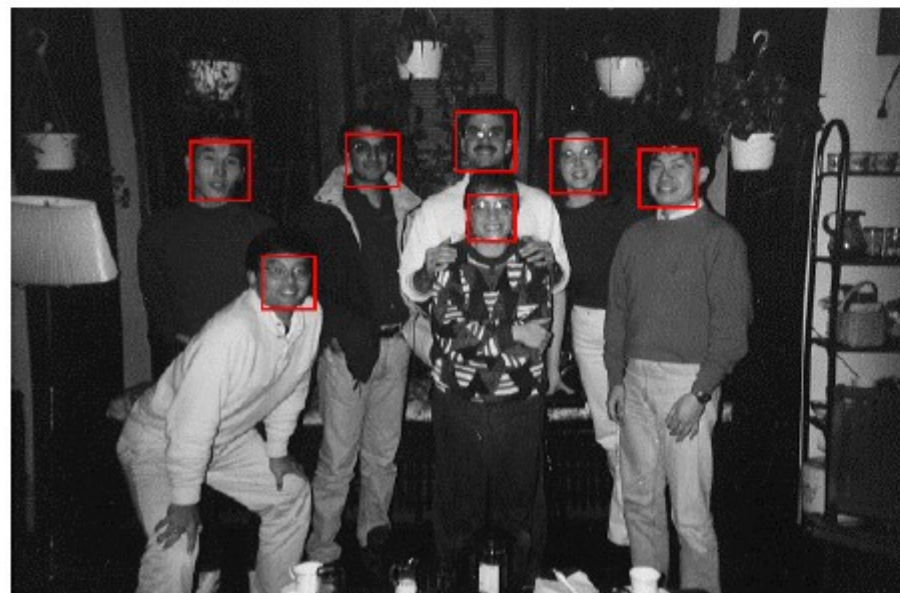
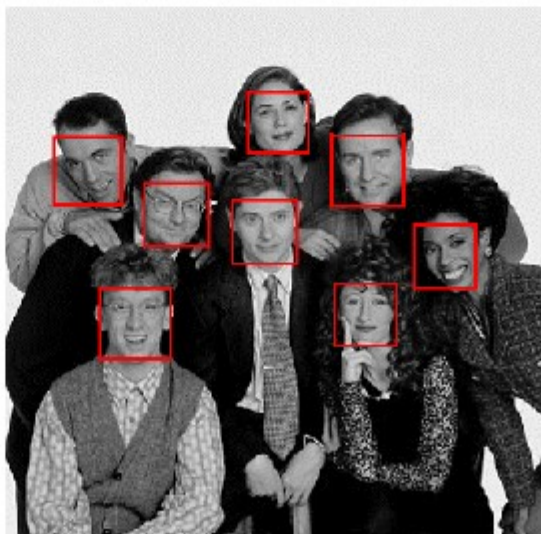
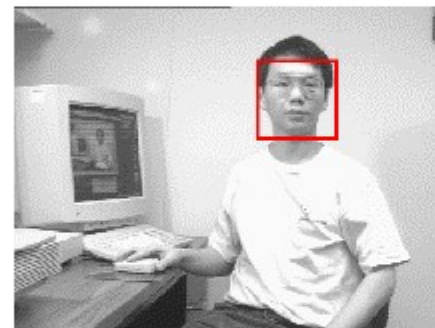
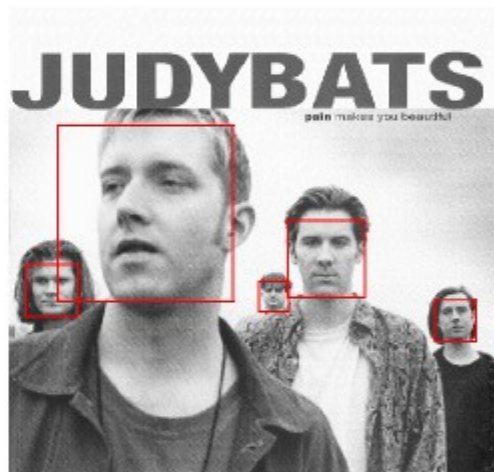
Performance evaluation –

# היום

- חזרה על סיווג – BOW + SVM
- זיהוי ע"י סיווג
- זיהוי פנים בשיטת Viola-jones
- Boosting
- Cascade classifier

# Detection via classification: Main idea

- **Consider all subwindows in an image**
  - Sample at multiple scales and positions
- **Make a decision per window:**
  - “Does this contain object category X or not?”
- **In this section, we'll focus specifically on methods using a global representation (i.e., not part-based, not local features).**



איתור פנים בשיטה של ויאולה  
וג'ונס

# Three Main Contributions

1. **Integral image**, rectangle features can be computed efficiently
2. **Boosting** - constructing a classifier using AdaBoost - learning algorithm developed by Freund and Schapire, selects a small set of features and build a classifier based on them
3. **Combining** successively more **complex classifiers** in a “**cascade**” – focus attention on promising regions of the picture

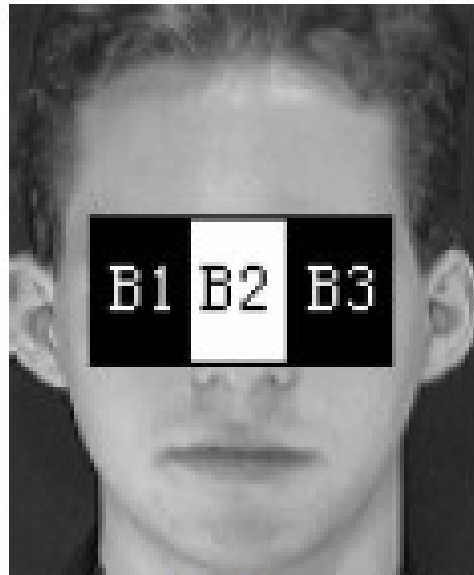
# בראשי פרקים

- Rectangle features
- Integral image
- Weak learner
- Strong learner
- The boosting algorithm
- The cascade

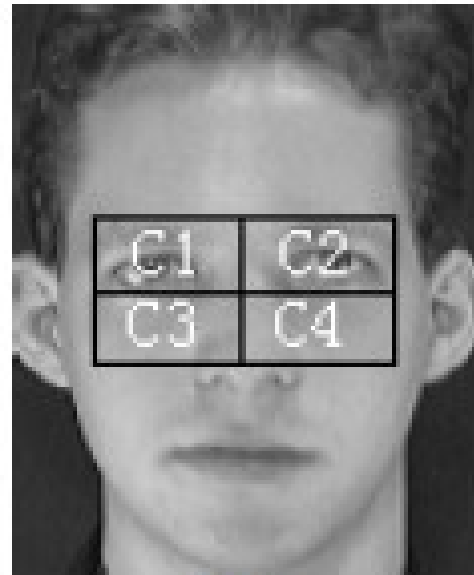
# Rectangle Features



(a)



(b)

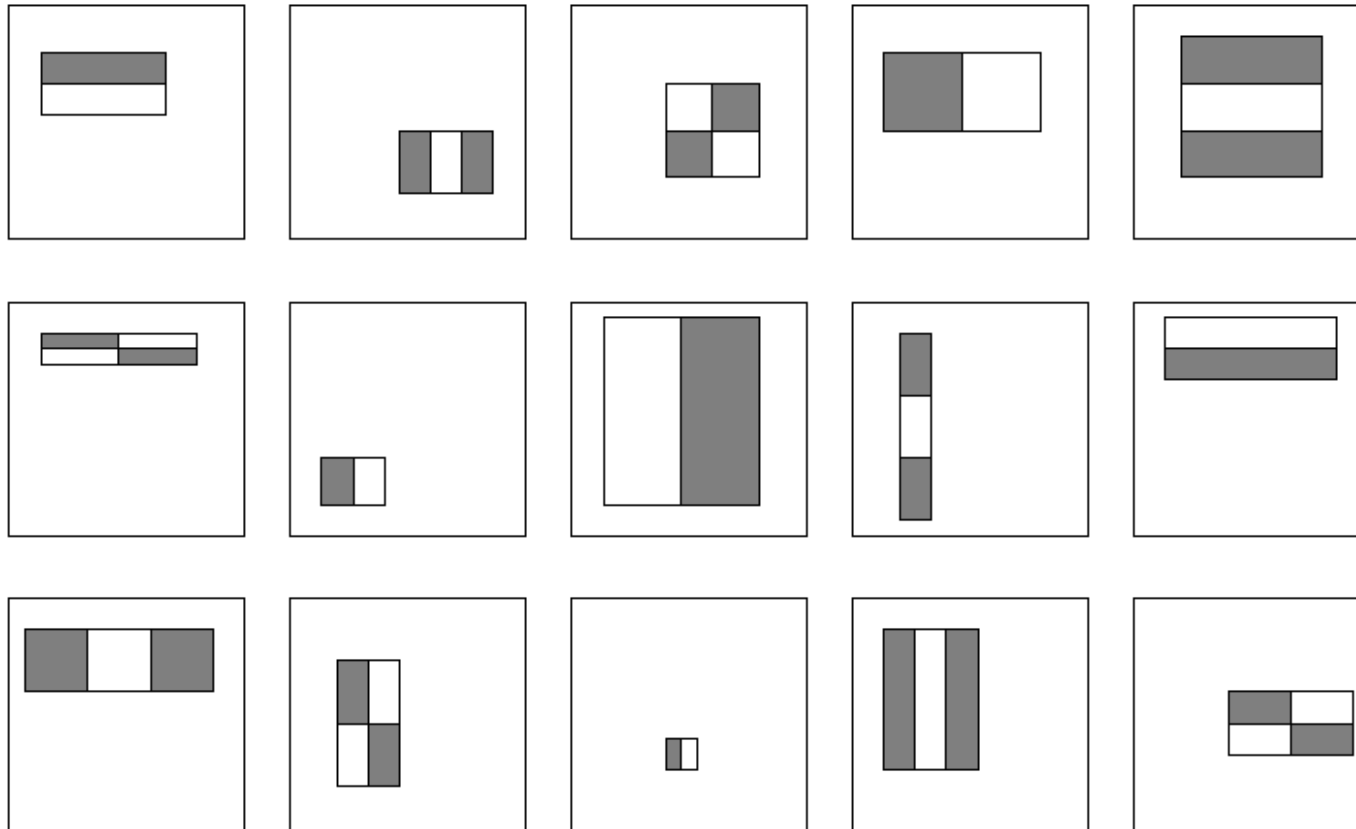


(c)



# Rectangle Features

For a 24x24 sub-window, the number of possible rectangle features is  $\sim 180,000$ !  
Select using AdaBoost...



# Integral Image Representation

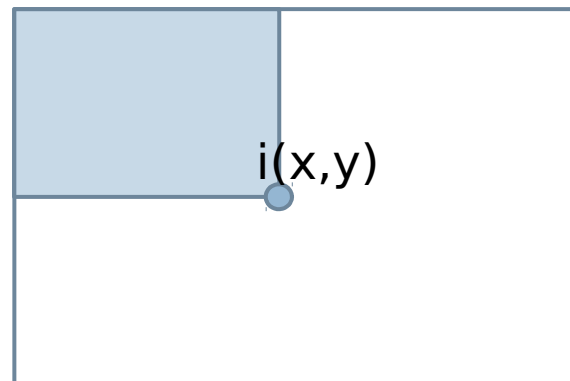
- Each pixel contains the **sum of the pixels above and to the left** of it:

$$ii(x, y) = \sum_{x' \leq x, y' \leq y} i(x', y')$$

The value in  
the integral  
image

The value in  
the original  
image

- Example:



# Boosting

- Build a strong classifier by combining number of “weak classifiers”, which need only be better than chance
- Sequential learning process: at each iteration, add a weak classifier
- Flexible to choice of weak learner
  - including fast simple classifiers that alone may be inaccurate
- We’ll look at Freund & Schapire’s AdaBoost algorithm
  - Easy to implement
  - Base learning algorithm for Viola-Jones face detector

# A Formal View of Boosting

- given training set  $(x_1, y_1), \dots, (x_m, y_m)$
- $y_i \in \{-1, +1\}$  correct label of instance  $x_i \in X$
- for  $t = 1, \dots, T$ :
  - construct distribution  $D_t$  on  $\{1, \dots, m\}$
  - find weak hypothesis (“rule of thumb”)  
 $h_t : X \rightarrow \{-1, +1\}$   
with small error  $\epsilon_t$  on  $D_t$ :  
 $\epsilon_t = \Pr_{D_t}[h_t(x_i) \neq y_i]$
- output final hypothesis  $H_{\text{final}}$

# AdaBoost

[Freund & Schapire]

- constructing  $D_t$ :

- $D_1(i) = 1/m$
- given  $D_t$  and  $h_t$ :

$$\begin{aligned} D_{t+1}(i) &= \frac{D_t(i)}{Z_t} \cdot \begin{cases} e^{-\alpha_t} & \text{if } y_i = h_t(x_i) \\ e^{\alpha_t} & \text{if } y_i \neq h_t(x_i) \end{cases} \\ &= \frac{D_t(i)}{Z_t} \cdot \exp(-\alpha_t y_i h_t(x_i)) \end{aligned}$$

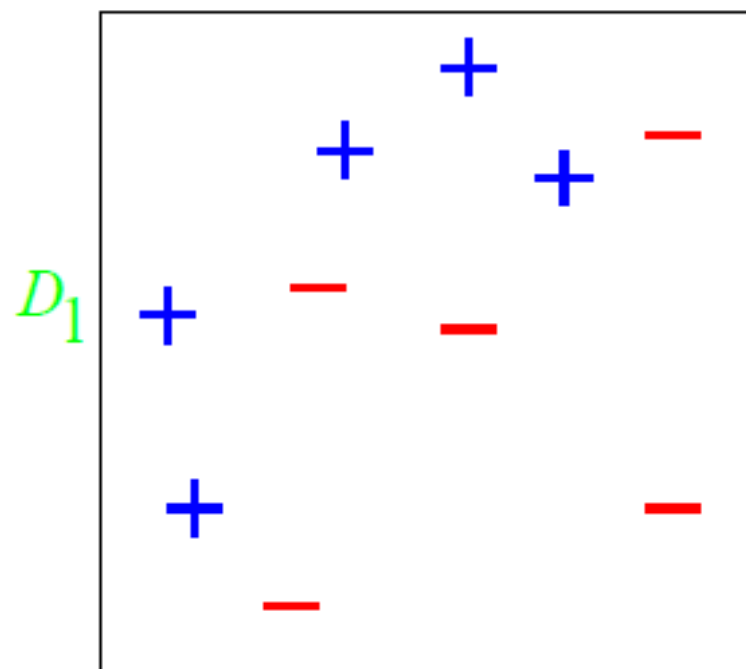
where  $Z_t = \text{normalization constant}$

$$\alpha_t = \frac{1}{2} \ln \left( \frac{1 - \epsilon_t}{\epsilon_t} \right) > 0$$

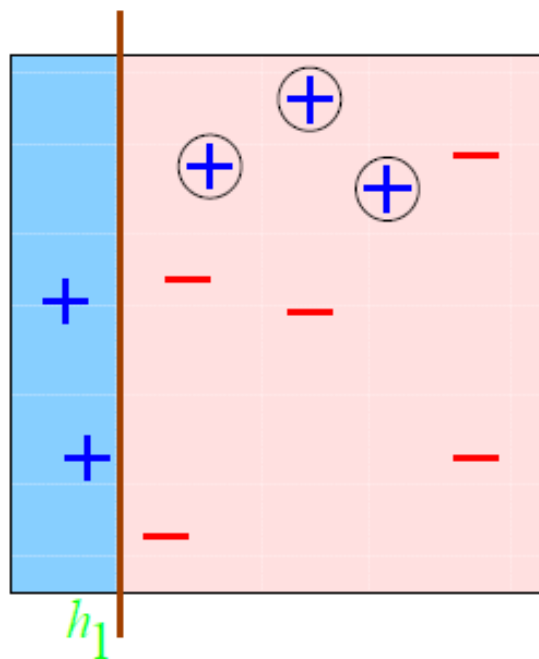
- final hypothesis:

- $H_{\text{final}}(x) = \text{sign} \left( \sum_t \alpha_t h_t(x) \right)$

## Toy Example

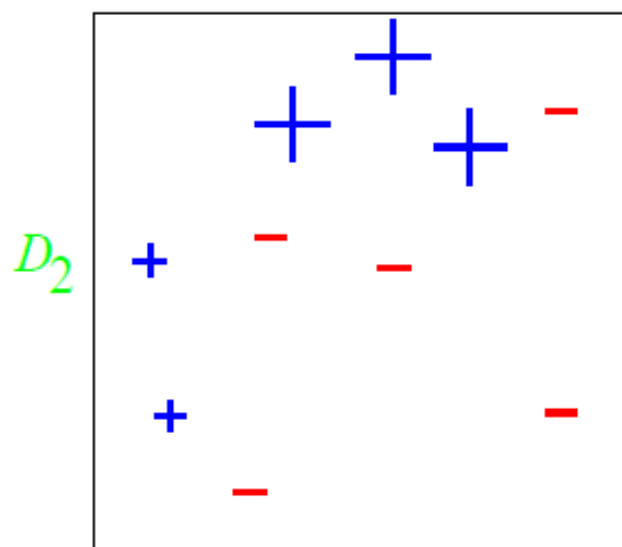


## Round 1

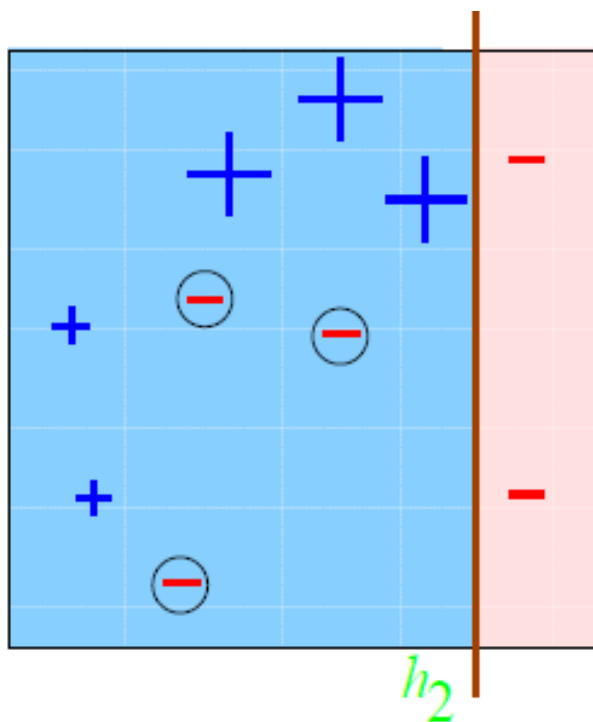
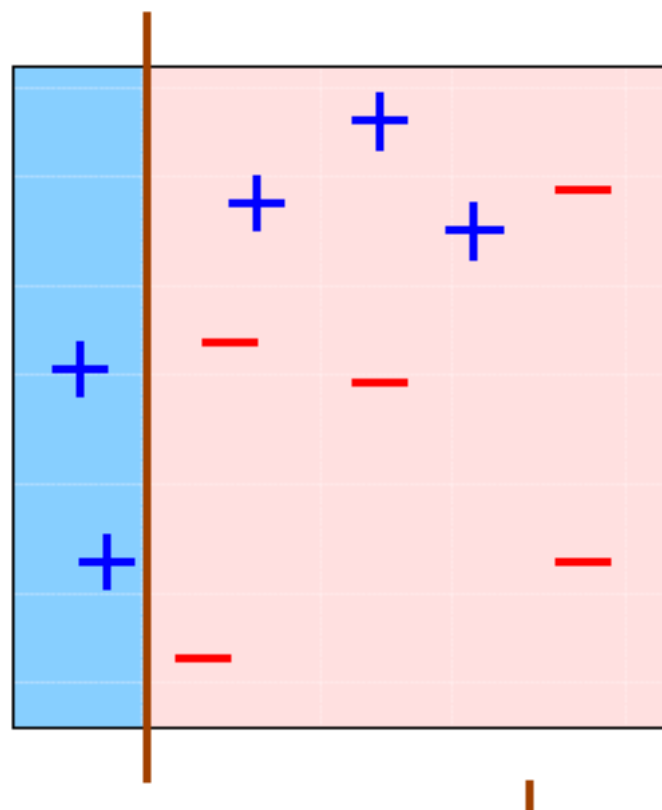


$$\epsilon_1 = 0.30$$

$$\alpha_1 = 0.42$$



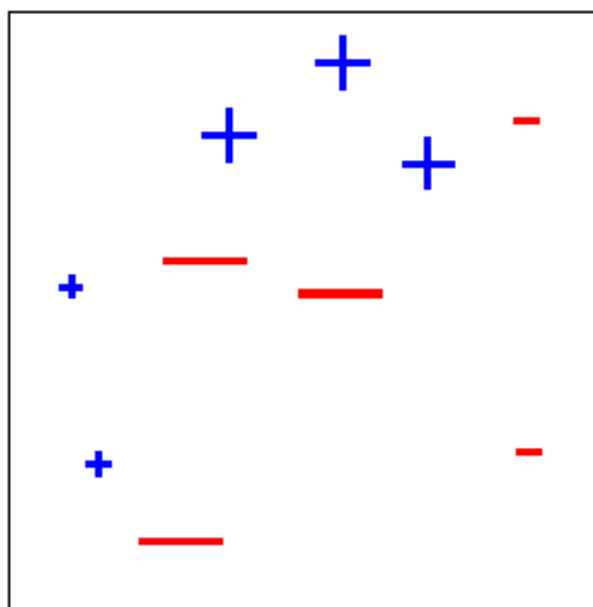
## Round 2



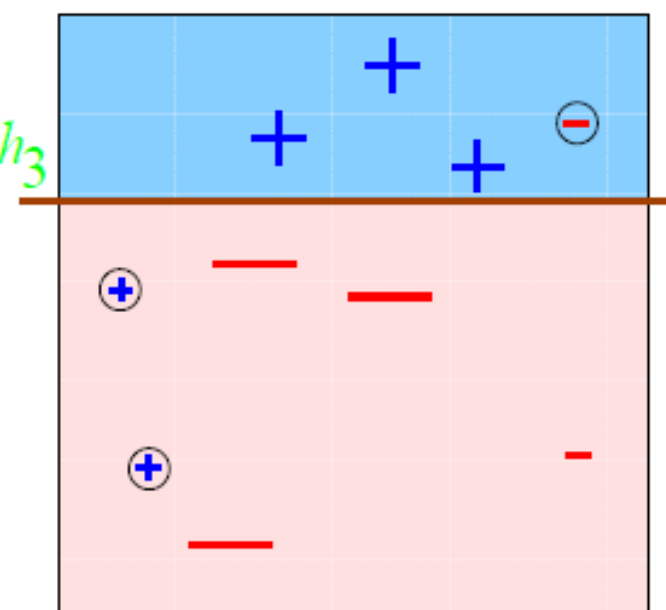
$$\epsilon_2 = 0.21$$

$$\alpha_2 = 0.65$$

$D_3$



$h_3$

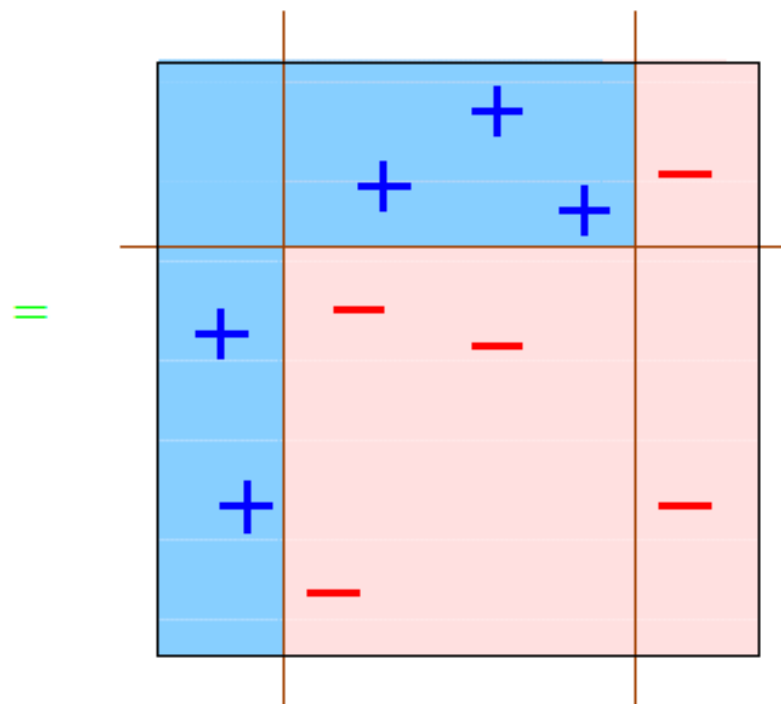




# Final Hypothesis

$H_{\text{final}}$

$$= \text{sign} \left( 0.42 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} + 0.65 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} + 0.92 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} \right)$$



- Given example images  $(x_1, y_1), \dots, (x_n, y_n)$  where  $y_i = 0, 1$  for negative and positive examples respectively.
- Initialize weights  $w_{1,i} = \frac{1}{2m}, \frac{1}{2l}$  for  $y_i = 0, 1$  respectively, where  $m$  and  $l$  are the number of negatives and positives respectively.
- For  $t = 1, \dots, T$ :

1. Normalize the weights,

$$w_{t,i} \leftarrow \frac{w_{t,i}}{\sum_{j=1}^n w_{t,j}}$$

so that  $w_t$  is a probability distribution.

- For each feature,  $j$ , train a classifier  $h_j$  which is restricted to using a single feature. The error is evaluated with respect to  $w_t$ ,  $\epsilon_j = \sum_i w_i |h_j(x_i) - y_i|$ .
- Choose the classifier,  $h_t$ , with the lowest error  $\epsilon_t$ .
- Update the weights:

$$w_{t+1,i} = w_{t,i} \beta_t^{1-e_i}$$

where  $e_i = 0$  if example  $x_i$  is classified correctly,  $e_i = 1$  otherwise, and  $\beta_t = \frac{\epsilon_t}{1-\epsilon_t}$ .

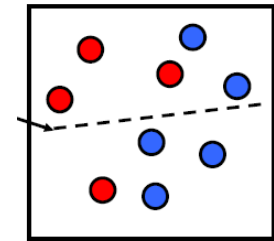
- The final strong classifier is:

$$h(x) = \begin{cases} 1 & \sum_{t=1}^T \alpha_t h_t(x) \geq \frac{1}{2} \sum_{t=1}^T \alpha_t \\ 0 & \text{otherwise} \end{cases}$$

where  $\alpha_t = \log \frac{1}{\beta_t}$

# AdaBoost Algorithm

Start with  
uniform weights  
on training  
examples



$\{x_1, \dots, x_n\}$

For  $T$  rounds

← Evaluate *weighted*  
error for each  
feature, pick best.

Re-weight the examples:

← Incorrectly classified -> more weight  
Correctly classified -> less weight

← Final classifier is combination of the weak  
ones, weighted according to error they  
had.

**Freund & Schapire 1995**

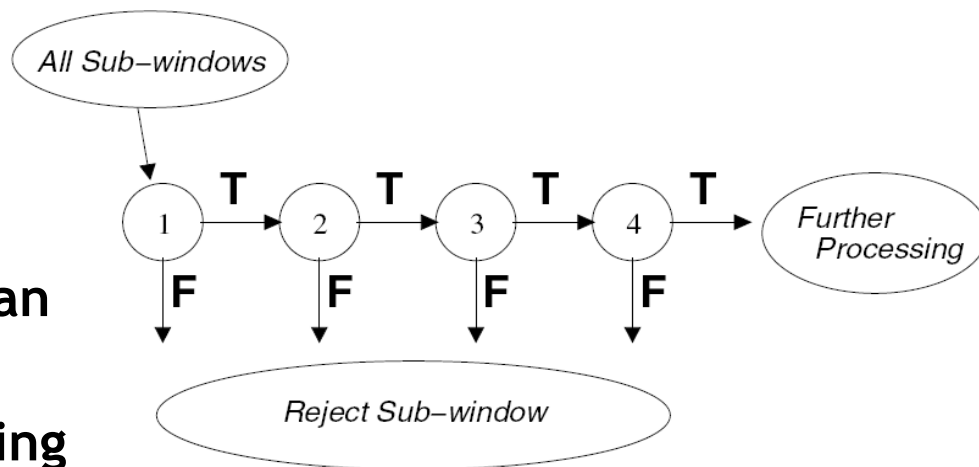
# Learning to Detect

- Image Features + thresholds = Weak Classifiers
- For each round of Boosting
  - Evaluate each rectangle filter on each example
  - Sort examples by filter values
  - Select best threshold for each filter (min. err.)
    - Sorted list can be quickly scanned for optimal threshold
  - Select best filter / threshold combo.
  - Weight on this feature is a simple function of err. rate.
  - Reweight examples
- Output: selected features and their weights

# Cascading classifiers for detection

For efficiency, apply less accurate but faster classifiers first to immediately discard windows that clearly appear to be negative; e.g.,

- Filter for promising regions with an initial inexpensive classifier
- Build a chain of classifiers, choosing cheap ones with low false negative rates early in the chain



Fleuret & Geman, IJCV 2001

Rowley et al., PAMI 1998

Viola & Jones, CVPR 2001

K. Grauman, B. Leibe

Figure from Viola & Jones CVPR 2001

9.8 % patches remaining



0.74 % patches remaining



0.06 % patches remaining

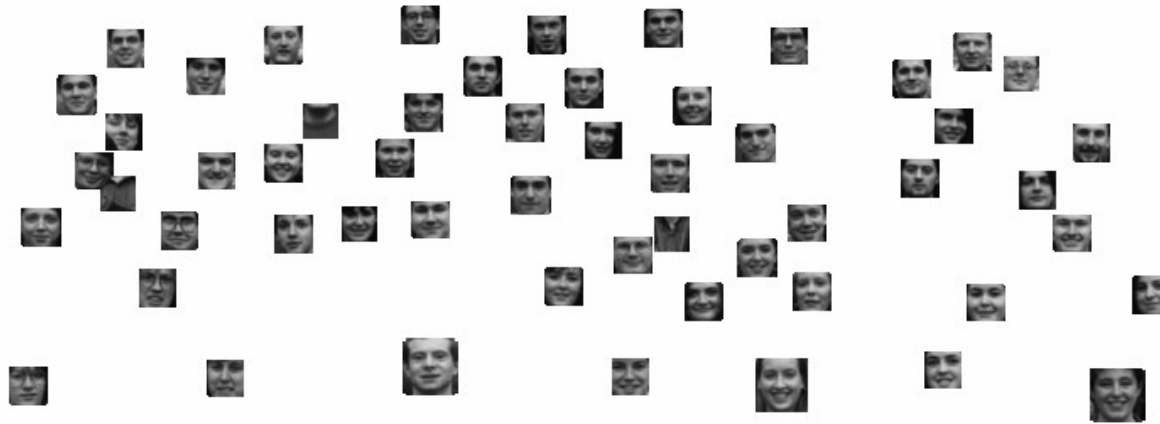


0.01 % patches remaining





0.007 % patches remaining



# Highlights

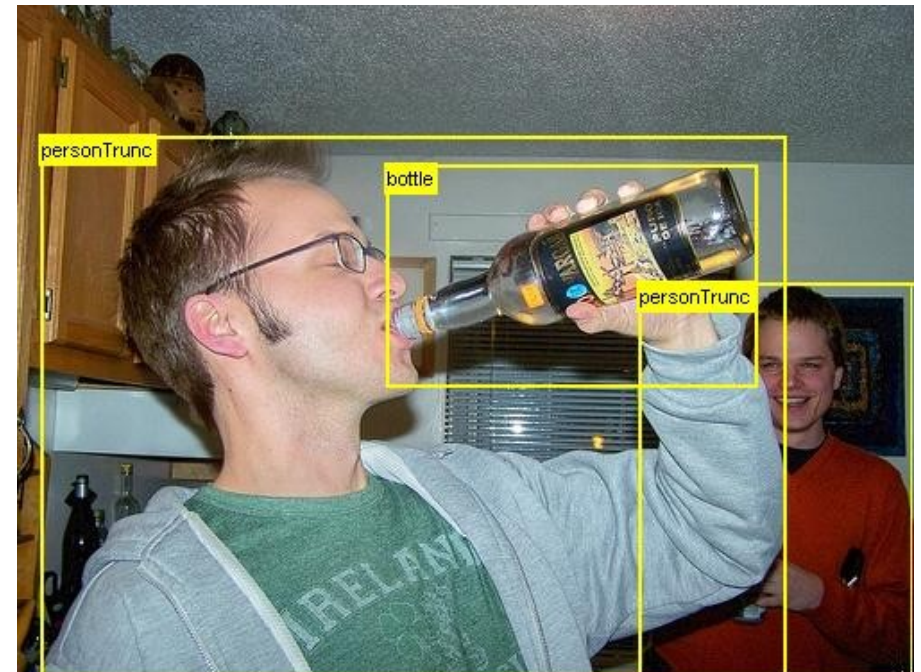
- **Sliding window detection and global appearance descriptors:**
  - Simple detection protocol to implement
  - Good feature choices critical
  - Past successes for certain classes

# Limitations

- **High computational complexity**
  - For example: 250,000 locations x 30 orientations x 4 scales = 30,000,000 evaluations!
  - If training binary detectors independently, means cost increases linearly with number of classes
- **With so many windows, false positive rate better be low**

# Limitations (continued)

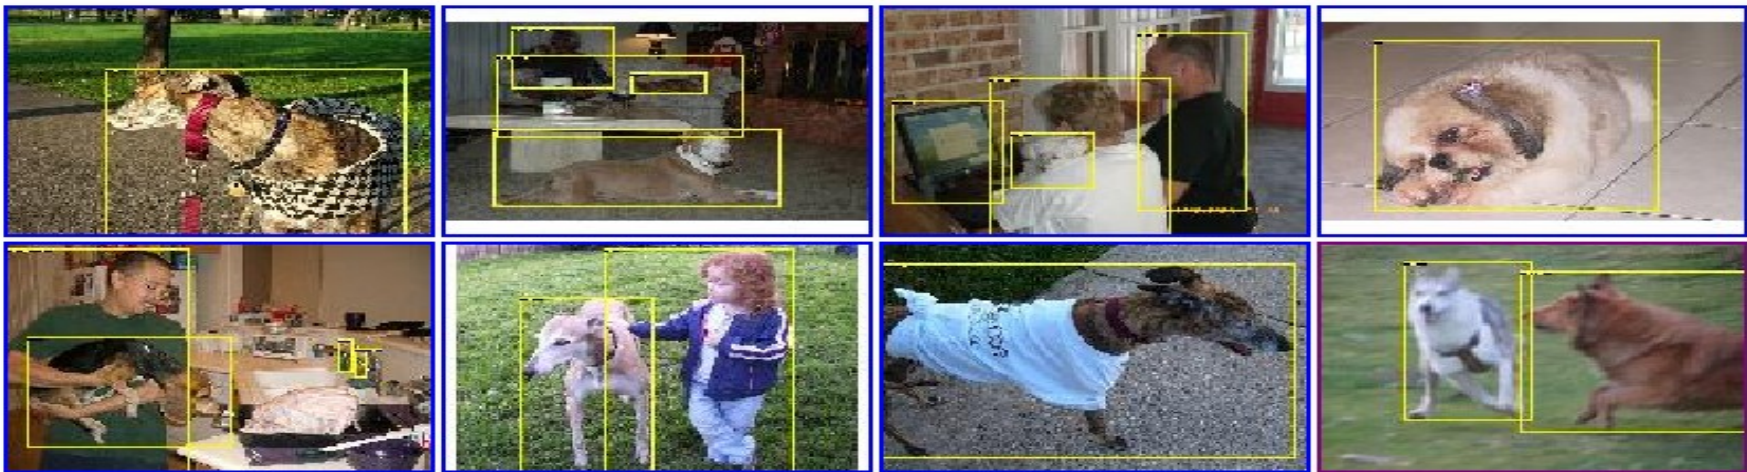
- Not all objects are “box” shaped



# Limitations (continued)

- Non-rigid, deformable objects not captured well with representations assuming a fixed 2d structure; or must assume fixed viewpoint
- Objects with less-regular textures not captured well with holistic appearance-based descriptions

*Dogs - all images contain at least one dog.*





# Limitations (continued)

- If considering windows in isolation, context is lost



**Sliding window**



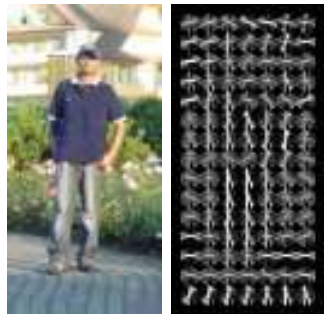
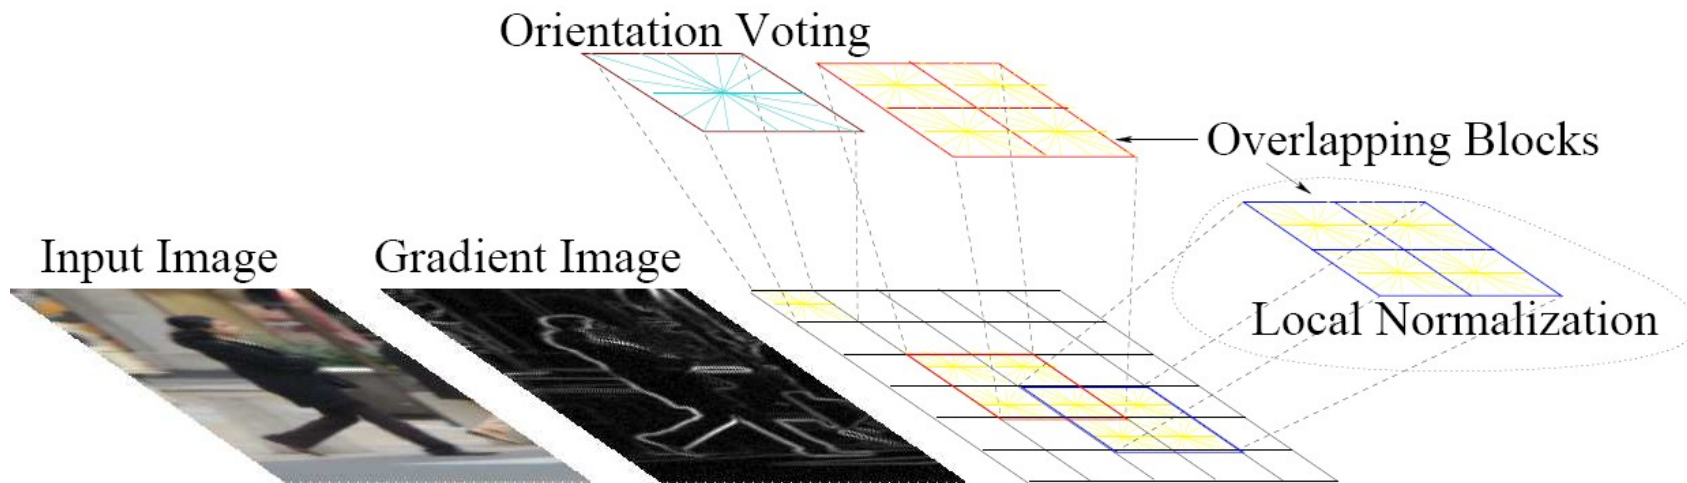
**Detector's view**

# Limitations (continued)

- In practice, often entails large, cropped training set (expensive)
- Requiring good match to a global appearance description can lead to sensitivity to partial occlusions



# Gradient-based representations: Histograms of oriented gradients (HoG)



**Map each grid cell in the input window to a histogram counting the gradients per orientation.**

**Code available:**  
<http://pascal.inrialpes.fr/soft/olt/>

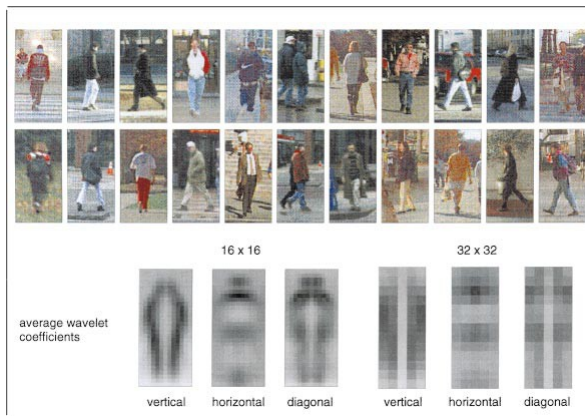
**Dalal & Triggs, CVPR 2005**

K. Grauman, B. Leibe



# Pedestrian detection

- **Detecting upright, walking humans also possible using sliding window's appearance/texture; e.g.,**



## SVM with Haar wavelets [Papageorgiou & Poggio, IJCV 2000]

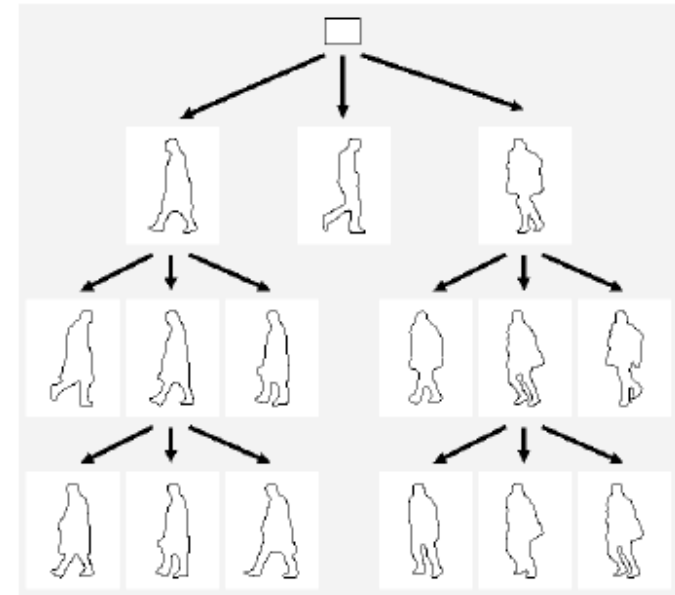
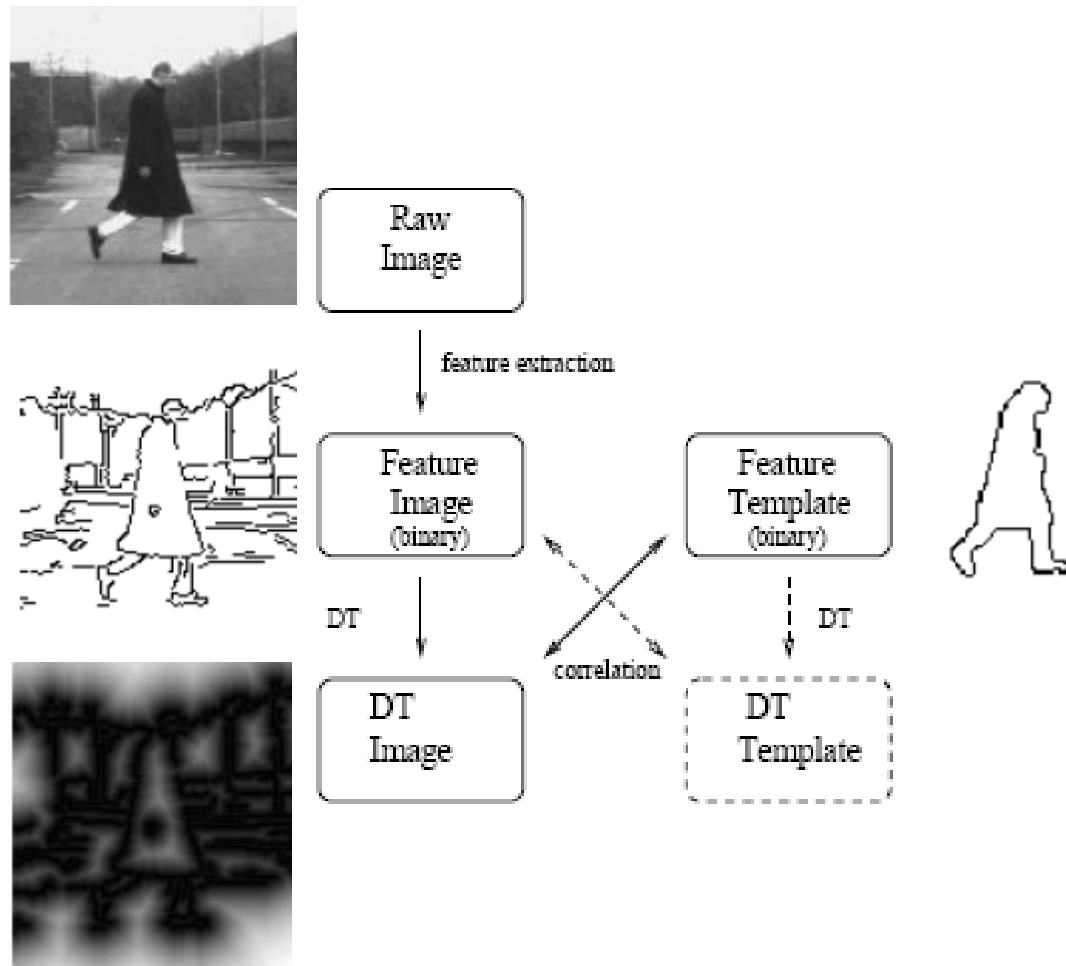


## Space-time rectangle features [Viola, Jones & Snow, ICCV 2003]



## SVM with HoGs [Dalal & Triggs, CVPR 2005]

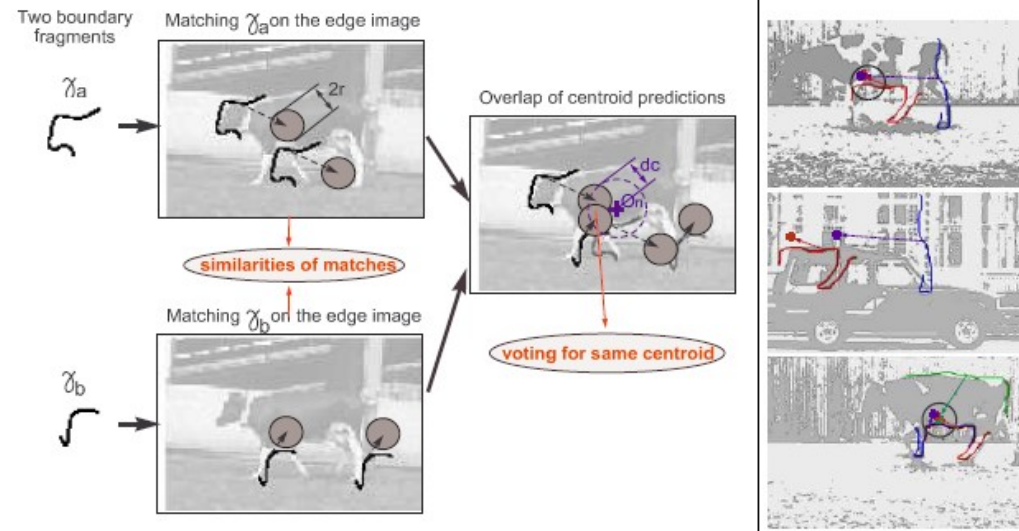
# Edges and chamfer distance



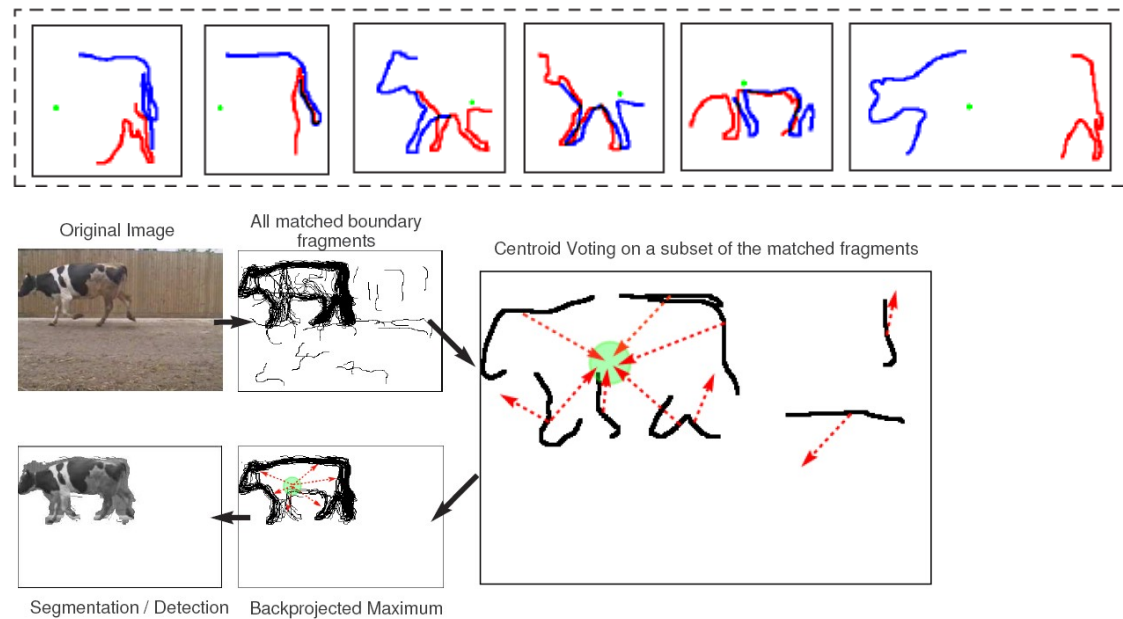
Gavrila, Philomin, ICCV 1999

# Edge fragments

Opelt, Pinz, Zisserman,  
ECCV 2006

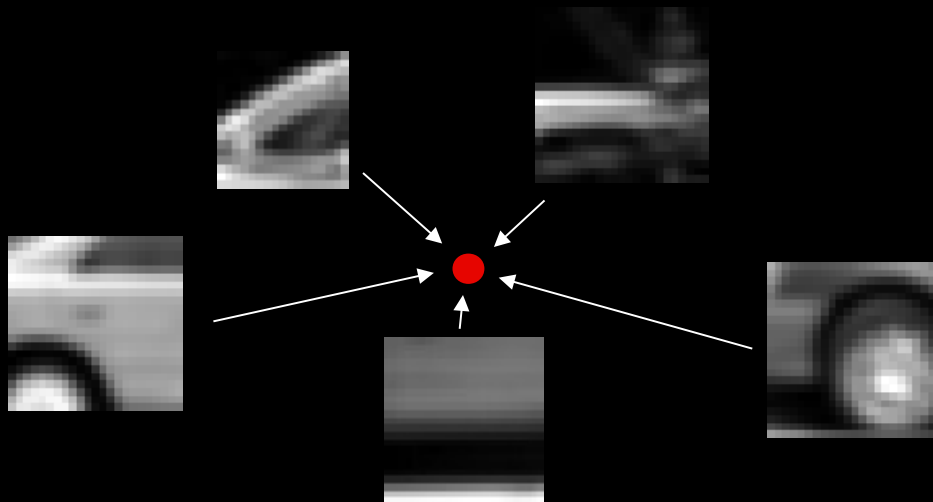


Weak detector = k edge  
fragments and threshold.  
Chamfer distance uses 8  
orientation planes

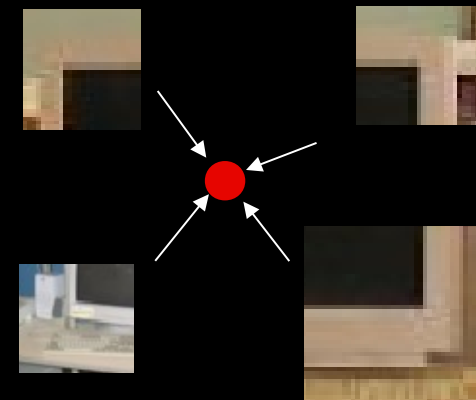


# Weak detectors

**Part based: similar to part-based generative models. We create weak detectors by using parts and voting for the object center location**



Car model



Screen model

These features are used for the detector on the course web site.

# סיכום

- זיהוי ע"י סיווג
- Boosting
- Viola & Jones

# בפעם הבאה

- סינטז טקסטורה
- גאומטריה של תמונות